

«МОСКОВСКИЙ ФИЗИКО-ТЕХНИЧЕСКИЙ ИНСТИТУТ (национальный
исследовательский университет)
ФИЗТЕХ-ШКОЛА ПРИКЛАДНОЙ МАТЕМАТИКИ И ИНФОРМАТИКИ
КАФЕДРА «ИНТЕЛЛЕКТУАЛЬНЫЕ СИСТЕМЫ»

Исламов Рустем Ильфакович

Распределенные методы второго порядка с быстрой скоростью сходимости и компрессией

03.03.01 — Прикладные математика и физика

ВЫПУСКНАЯ КВАЛИФИКАЦИОННАЯ РАБОТА БАКАЛАВРА

Научный руководитель:

д. ф.-м. н. Стрижов Вадим Викторович

Москва
2021

Аннотация

Данная бакалаврская диссертация основана на статье «Distributed Second Order Methods with Fast Rates and Compressed Communication» [17] за авторством Рустема Исламов, Шун Кяна и Питера Рихтарика.

Мы разработали несколько новых эффективных с точки зрения коммуникации методов второго порядка для распределенной оптимизации. Наш первый метод, NEWTON-STAR, является одним из вариантов метода Ньютона, от которого он наследует свою быструю локальную квадратичную скорость. Однако, в отличие от метода Ньютона, NEWTON-STAR имеет ту же стоимость коммуницирования за одну итерацию, что и градиентный спуск. Хотя этот метод непрacticен, поскольку он опирается на использование некоторых неизвестных параметров, характеризующих Гессиан целевой функции в оптимуме, он служит отправной точкой, которая позволяет нам проектировать практические варианты с доказанными теоретическими гарантиями сходимости. В частности, мы разрабатываем стратегию, основанную на использовании случайной разреженности для изучения неизвестных параметров итеративным способом, сохраняя эффективность с точки зрения коммуникаций. Применение этой стратегии к NEWTON-STAR приводит к нашему следующему методу, NEWTON-LEARN, для которого мы доказали локальные линейные и сверхлинейные скорости сходимости, не зависящие от числа обусловленности функции. Когда эти методы применимы, они имеют значительно более высокие скорости сходимости по сравнению с современными методами. Наконец, мы разработали стратегию глобализации с использованием кубической регуляризации, которая приводит к нашему следующему методу, CUBIC-NEWTON-LEARN, для которого мы доказали глобальные сублинейные и линейные скорости сходимости, а также быструю сверхлинейную скорость сходимости. Наши результаты подтверждаются экспериментальными результатами на реальных наборах данных и показывают улучшение на несколько порядков по сравнению с базовыми и современными методами с точки зрения эффективности коммуницирования.

Abstract

This bachelor thesis is based on paper “Distributed Second Order Methods with Fast Rates and Compressed Communication” [17] written by Rustem Islamov, Xun Qian, and Peter Richtárik.

We develop several new communication-efficient second-order methods for distributed optimization. Our first method, **NEWTON-STAR**, is a variant of Newton’s method from which it inherits its fast local quadratic rate. However, unlike Newton’s method, **NEWTON-STAR** enjoys the same per iteration communication cost as gradient descent. While this method is impractical as it relies on the use of certain unknown parameters characterizing the Hessian of the objective function at the optimum, it serves as the starting point which enables us design practical variants thereof with strong theoretical guarantees. In particular, we design a stochastic sparsification strategy for learning the unknown parameters in an iterative fashion in a communication efficient manner. Applying this strategy to **NEWTON-STAR** leads to our next method, **NEWTON-LEARN**, for which we prove local linear and superlinear rates independent of the condition number. When applicable, this method can have dramatically superior convergence behavior when compared to state-of-the-art methods. Finally, we develop a globalization strategy using cubic regularization which leads to our next method, **CUBIC-NEWTON-LEARN**, for which we prove global sublinear and linear convergence rates, and a fast superlinear rate. Our results are supported with experimental results on real datasets, and show several orders of magnitude improvement on baseline and state-of-the-art methods in terms of communication complexity.

Contents

1	Introduction	6
1.1	Distributed optimization	6
1.2	The curse of the condition number	7
1.3	Newton’s method to the rescue?	7
1.4	Contributions summary	8
1.5	Related work	9
2	Three Steps Towards an Efficient Distributed Newton Type Method	11
2.1	Naive distributed implementation of Newton’s method	11
2.2	A better implementation taking advantage of the structure of $\mathbf{H}_{ij}(x)$. . .	12
2.3	NEWTON-STAR: Newton’s method with a single Hessian	12
3	NEWTON-LEARN: Learning the Hessian and Local Convergence Theory	14
3.1	The main iteration	14
3.2	Learning the coefficients: the idea	15
3.3	Outline of fast local convergence theory	15
3.4	Compressed learning	16
3.5	NL1 (learning in the $\lambda > 0$ case)	16
3.5.1	The learning iteration and the NL1 algorithm	17
3.5.2	Theory	17
3.6	NL2 (learning in the $\lambda \geq 0$ case)	18
3.6.1	The learning iteration and the NL2 algorithm	18
3.6.2	Theory	19
4	CUBIC-NEWTON-LEARN: Global Convergence Theory via Cubic Regularization	20
4.1	CNL: the algorithm	21
4.2	Global convergence	22
4.3	Superlinear convergence	23
5	Experiments	23
5.1	Data sets and parameter settings	24
5.2	Compression operators	24
5.2.1	Random sparsification	25
5.2.2	Random dithering	25
5.2.3	Natural compression	25
5.2.4	Bernoulli compressor	25
5.3	Behavior of NL1 and NL2	26
5.4	Comparison of NL1 and NL2 with Newton’s method	26
5.5	Comparison of NL1 and NL2 with BFGS	26
5.6	Comparison of NL1 and NL2 with ADIANA	27
5.7	Comparison of NL1 and NL2 with DINGO	27
5.8	Comparison of CNL with DCGD and DIANA	28
A	Proofs for NL1 (Section 3.5)	30
A.1	Lemma	30

A.2	Proof of Theorem 3.2	31
A.3	Proof of Lemma 3.3	34
B	Proofs for NL2 (Section 3.6)	36
B.1	Proof of Theorem 3.5	36
B.2	Proof of Lemma 3.6	39
C	Proofs for CNL (Section 4)	40
C.1	Solving the Subproblem	40
C.2	Proof of Lemma 4.2	40
C.3	Proof of Theorem 4.3	41
C.4	Proof of Theorem 4.4	42
C.5	Proof of Theorem 4.5	43
D	Extra Method: MAX-NEWTON	45
	References	48

1 Introduction

The prevalent paradigm for training modern supervised machine learning models is based on (regularized) empirical risk minimization (ERM) [41], and the most commonly used optimization methods deployed for solving ERM problems belong to the class of stochastic first order methods [32, 38]. Since modern training data sets are very large and are becoming larger every year, it is increasingly harder to get by without relying on modern computing architectures which make efficient use of distributed computing. However, in order to develop efficient distributed methods, one has to keep in mind that communication among the different parallel workers (e.g. processors or compute nodes) is typically very slow, and almost invariably forms the main bottleneck in deployed optimization software and systems [3]. For this reason, further advances in the area of communication efficient distributed first order optimization methods for solving ERM problems are highly needed, and research in this area constitutes one of the most important fundamental endeavors in modern machine learning. Indeed, this research field is very active, and numerous advances have been made over the past decade [1, 4, 29, 40, 44, 48, 51].

1.1 Distributed optimization

We consider L2 regularized empirical risk minimization problems of the form

$$\min_{x \in \mathbb{R}^d} \left[P(x) := f(x) + \frac{\lambda}{2} \|x\|^2 \right], \quad (1)$$

where $f : \mathbb{R}^d \rightarrow \mathbb{R}$ is a smooth¹ convex function of the “average of averages” structure

$$f(x) := \frac{1}{n} \sum_{i=1}^n f_i(x), \quad f_i(x) := \frac{1}{m} \sum_{j=1}^m f_{ij}(x), \quad (2)$$

and $\lambda \geq 0$ is a regularization parameter. Here n is the number of parallel workers (nodes), and m is the number of training examples handled by each node². The value $f_{ij}(x)$ denotes the loss of the model parameterized by vector $x \in \mathbb{R}^d$ on the j^{th} example owned by the i^{th} node. This example is denoted as $a_{ij} \in \mathbb{R}^d$, and the corresponding loss function is $\varphi_{ij} : \mathbb{R} \rightarrow \mathbb{R}$, and hence we have

$$f_{ij}(x) := \varphi_{ij}(a_{ij}^\top x). \quad (3)$$

Thus, f represents the average loss/risk over all nm training datapoints, and problem (1) seeks to find the model whose (L2 regularized) empirical risk is minimized. We make the following assumption throughout the paper.

Assumption 1.1. *Problem (1) has at least one optimal solution x^* . For all i and j , the loss function $\varphi_{ij} : \mathbb{R} \rightarrow \mathbb{R}$ is γ -smooth, twice differentiable, and its second derivative $\varphi_{ij}'' : \mathbb{R} \rightarrow \mathbb{R}$ is ν -Lipschitz continuous.*

¹Function $\phi : \mathbb{R}^d \rightarrow \mathbb{R}$ is *smooth* if it is differentiable, and has L_ϕ Lipschitz gradient: $\|\nabla\phi(x) - \nabla\phi(y)\| \leq L_\phi \|x - y\|$ for all $x, y \in \mathbb{R}^d$. We say that L_ϕ is the *smoothness constant* of ϕ .

²All our results can be extended in a straightforward way to the more general case when node i contains m_i training examples. We decided to present the results in the special case $m = m_i$ for all i in order to simplify the notation.

Note that in view of (3), the Hessian of f_{ij} at point x is

$$\mathbf{H}_{ij}(x) := \nabla^2 f_{ij}(x) = h_{ij}(x) a_{ij} a_{ij}^\top, \quad (4)$$

where

$$h_{ij}(x) := \varphi''_{ij}(a_{ij}^\top x). \quad (5)$$

In view of Assumption 1.1, we have $|\varphi''_{ij}(t)| \leq \gamma$ for all $t \in \mathbb{R}$, and

$$|h_{ij}(x) - h_{ij}(y)| \leq \nu |a_{ij}^\top x - a_{ij}^\top y| \leq \nu \|a_{ij}\| \|x - y\| \quad (6)$$

for all $x, y \in \mathbb{R}^d$. Let $R := \max_{ij} \|a_{ij}\|$. The Hessian of f_i is given by

$$\mathbf{H}_i(x) \stackrel{(2)}{=} \frac{1}{m} \sum_{j=1}^m \mathbf{H}_{ij}(x) \stackrel{(4)}{=} \frac{1}{m} \sum_{j=1}^m h_{ij}(x) a_{ij} a_{ij}^\top, \quad (7)$$

and the Hessian of f is given by

$$\mathbf{H}(x) \stackrel{(2)}{=} \frac{1}{n} \sum_{i=1}^n \mathbf{H}_i(x) \stackrel{(7)}{=} \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m h_{ij}(x) a_{ij} a_{ij}^\top. \quad (8)$$

1.2 The curse of the condition number

All first order methods—distributed or not—suffer from a dependence on an appropriately chosen notion of a *condition number*³—a number that describes the difficulty of solving the problem by the method at hand. A condition number is a function of the goal we are trying to achieve (e.g., minimize the number of iterations vs minimize the number of communications), choice of the loss function, structure of the model we are trying to learn, and last but not least, the size and properties of the training data. In fact, most research in this area is motivated by the desire to design methods that would have a *reduced* dependence on the condition number. This is the case for many of the tricks heavily studied in the literature, including minibatching [47], importance sampling [31, 55], random reshuffling [30], variance reduction [8, 18, 39, 52], momentum [25, 26], adaptivity [28], communication compression [1, 4, 29], and local computation [19, 27, 45]. Research in this area is becoming saturated, and new ideas are needed to make further progress.

1.3 Newton’s method to the rescue?

One of the ideas that undoubtedly crossed everybody’s mind is the trivial observation that there *is* a very old and simple method which does *not* suffer from any conditioning issues: Newton’s method. Indeed, when it works, Newton’s method has a fast *local quadratic convergence rate* which is entirely independent of the condition number of the problem [2]. While this is a very attractive property, developing scalable distributed variants of Newton’s method that could also *provably outperform* gradient based methods remains a largely unsolved problem. To highlight the severity of the issues with extending Newton’s method to stochastic and distributed settings common in machine learning, we note that

³Example: if one wishes to minimize an L -smooth μ -strongly convex function and one cares about the number of gradient type iterations, the appropriate notion of a condition number is $\kappa := \frac{L}{\mu}$.

Table 1: Summary of algorithms proposed and convergence results proved in this paper.

Method	Convergence			Rate of the condition number?	Theorem
	result [†]	type	rate		
NEWTON-STAR (12)	$r_{k+1} \leq cr_k^2$	local	quadratic	✓	2.1
MAX-NEWTON Algorithm 4	$r_{k+1} \leq cr_k^2$	local	quadratic	✓	D.1
NEWTON-LEARN Algorithm 1	$\Phi_1^k \leq \theta_1^k \Phi_1^0$	local	linear	✓	3.2
	$r_{k+1} \leq c\theta_1^k r_k$	local	superlinear	✓	3.2
NEWTON-LEARN Algorithm 2	$\Phi_2^k \leq \theta_2^k \Phi_2^0$	local	linear	✓	3.5
	$r_{k+1} \leq c\theta_2^k r_k$	local	superlinear	✓	3.5
CUBIC-NEWTON-LEARN Algorithm 3	$\Delta_k \leq \frac{\epsilon}{k}$	global	sublinear	✗	4.3
	$\Delta_k \leq c \exp(-k/c)$	global	linear	✗	4.4
	$\Phi_3^k \leq \theta_3^k \Phi_3^0$	local	linear	✓	4.5
	$r_{k+1} \leq c\theta_3^k r_k$	local	superlinear	✓	4.5

Quantities for which we prove convergence: (i) distance to solution $r_k := \|x^k - x^*\|$; (ii) Lyapunov function $\Phi_q^k := \|x^k - x^*\|^2 + c_q \sum_{i=1}^n \sum_{j=1}^m (h_{ij}^k - h_{ij}(x^*))^2$ for $q = 1, 2, 3$, where $h_{ij}(x^*) = \varphi_{ij}''(a_{ij}^\top x^*)$ (see (5)); (iii) Function value suboptimality $\Delta_k := P(x^k) - P(x^*)$

[†] constant c is possibly different each time it appears in this table. Refer to the precise statements of the theorems for the exact values.

until recently, we did not even have any Newton-type analogue of SGD that could provably work with small minibatch sizes, let alone minibatch size one [21]. In contrast, SGD with minibatch size one is one of the simplest and well understood variants thereof [31], and much of modern development in the area of SGD methods is much more sophisticated. Most variants of Newton’s method proposed for deployment in machine learning are heuristics, which is to say that they are not supported with any convergence guarantees, or have convergence guarantees without explicit rates, or suffer from rates that are worse than the rates of first order methods.

1.4 Contributions summary

We develop *several new fundamental Newton-type methods* which we hope make a marked step towards the ultimate goal of developing practically useful and *communication efficient distributed second order methods*. Our methods are designed with the explicit goal of supporting *efficient communication* in a distributed setting, and in sharp contrast with most recent work, their design was heavily influenced by our desire to equip them with *strong convergence guarantees* typical for the classical Newton’s method [36, 49] and cubically regularized Newton’s method [13, 33]. Our convergence results are summarized in Table 1.

- **First new method and its local quadratic convergence.** We first show that if we know the Hessian of the objective function at the optimal solution, then we can use it instead of the typical Hessian appearing in Newton’s method, and the resulting algorithm, which we call **NEWTON-STAR (NS)**, inherits local quadratic convergence behavior of Newton’s method (see Theorem 2.1). In a distributed setting with a central orchestrating sever, each compute node only needs to send the local gradient to the server node, and no matrices need to be sent. While this method is not practically useful, it acts as a stepping stone to our next method, in which

these deficiencies are removed. This method is described in Section 2. A somewhat different method with similar properties, which we call MAX-NEWTON⁴, is described in Section D.

- Second new method and its local linear and superlinear convergence.** Motivated by the above result, we propose a *learning scheme which enables us to learn the Hessian at the optimum iteratively in a communication efficient manner*. This scheme gives rise to our second new method: **NEWTON-LEARN (NL)**. We analyze this method in two cases: (i) all individual loss functions are convex and $\lambda > 0$ (giving rise to the NL1 method), and (ii) the aggregate loss function P is strongly convex (giving rise to the NL2 method). Besides the local full gradient, each worker node needs to send additional information to the server node in order to learn the Hessian at the optimum. However, our learning scheme supports *compressed communication* with arbitrary compression level. This level can be chosen so that in each iteration, each node sends an equivalent of a few gradients to the server only. That is, we can achieve $O(d)$ communication complexity in each iteration. In both cases, we prove local linear convergence for a carefully designed Lyapunov function, and local superlinear convergence for the squared distance to optimum (see Theorems 3.2 and 3.5). Remarkably, all these rates are *independent of the condition number*. The NL1 and NL2 methods and the associated theory are described in Section 3.
- Third new method and its global convergence.** Next, we equip our learning scheme with a cubic regularization strategy [13, 33], which leads to a new *globally* convergent method: **CUBIC-NEWTON-LEARN (CNL)**. We establish global sublinear and linear convergence (for function values) guarantees for convex and strongly convex problems, respectively. The method can also achieve a fast local linear (for a Lyapunov function) and superlinear (for squared distance to solution) convergence in the strongly convex case. We describe this method and the associated theory in Section 4.
- Experiments.** Our theory is corroborated with numerical experiments showing the superiority of our methods to several state-of-the-art benchmarks, including DCGD [20], DIANA [16, 29], ADIANA [23], BFGS [6, 9, 11, 43], and DINGO [7]. Our methods can achieve communication complexity which is *several orders of magnitude better* than competing methods (see Section 5).

1.5 Related work

Several distributed Newton-type methods can be found in recent literature. DANE [42] is a distributed approximate Newton-type method where each worker node needs to solve a subproblem using the full gradient at each iteration, and the new iterate is the average of these subproblem solutions. The linear convergence of DANE was obtained in the strongly convex case. An inexact DANE method in which the subproblem is solved approximately was proposed and studied by Reddi et al. [37]. Moreover, an accelerated version of inexact

⁴In fact, this was the first method we developed, in Summer 2020, when we embarked on the research which eventually lead to the results presented in this paper.

Table 2: Comparison of distributed Newton-type methods. Our methods combine the best of both worlds, and are the only methods we know about which do so: we obtain fast rates independent of the condition number, and allow for $O(d)$ communication per communication round.

Method	Convergence rate	Rate independent of the condition number?	Communication cost per iteration	Network structure
DANE [42]	Linear	✗	$O(d)$	Centralized
DiSCO [54]	Linear	✗	$O(d)$	Centralized
AIDE [37]	Linear	✗	$O(d)$	Centralized
GIANT [50]	Linear	✗	$O(d)$	Centralized
DINGO [7]	Linear	✗	$O(d)$	Centralized
DAN [53]	Local quadratic [†]	✓	$O(nd^2)$	Decentralized
DAN-LA [53]	Superlinear	✓	$O(nd)$	Decentralized
NEWTON-STAR this work	Local quadratic	✓	$O(d)$	Centralized
MAX-NEWTON this work	Local quadratic	✓	$O(d)$	Centralized
NEWTON-LEARN this work	Local superlinear	✓	$O(d)$	Centralized
CUBIC-NEWTON-LEARN this work	Superlinear	✓	$O(d)$	Centralized

[†] DAN converges globally, but the quadratic rate is introduced only after $O(L_2/\mu^2)$ steps, where L_2 is the Lipschitz constant of the Hessian of P , and μ is the strong convexity parameter of P . This is a property it inherits from the recent method of Polyak [34] this method is based on.

DANE, called AIDE, was proposed in [37] by a generic acceleration scheme—catalyst [24]—and an optimal communication complexity can be obtained up to logarithmic factors in specific settings. The DiSCO method, which combines inexact damped Newton method and distributed preconditioned conjugate gradient method, was proposed by Zhang and Xiao [54] and analyzed for self-concordant empirical loss. GIANT [50] is a globally improved approximate Newton method which has a better linear convergence rate than first-order methods for quadratic functions, and has local linear-quadratic convergence for strongly convex functions. GIANT and DANE are identical for quadratic programming. The communication cost per iteration of the above methods is $O(d)$. These methods can only achieve linear convergence in the strongly convex case. The comparison of the iteration complexity of the above methods for the ridge regression problem can be found in Table 2 of [50].

Crane and Roosta [7] proposed a distributed Newton-type method called DINGO for solving invex finite-sum problems. Invexity is a special case of non-convexity, which subsumes convexity as a sub-class. A linear convergence rate was obtained for DINGO under certain assumptions using an Armijo-type line search, and at each iteration, several communication rounds are needed assuming two communication rounds for line-search per iteration. The communication cost for each communication round is $O(d)$. The compressed version of DINGO was studied in [10] to reduce the communication cost

at each communication round by using the δ -approximate compressor, and the same rate of convergence as DINGO can be obtained by properly choosing the stepsize and hyper-parameters when δ is a constant. Zhang et al. [53] proposed two decentralized distributed adaptive Newton methods, called DAN and DAN-LA. DAN combines the distributed selective flooding (DSF) algorithm and Polyak’s adaptive Newton method [35], and enters pure Newton method which has quadratic convergence after about $\frac{2M}{\mu^2} \|\nabla P(x^0)\|$ iterations, where M is the Lipschitz constant of the Hessian of P and μ is the strongly convex parameter of P . DAN-LA, which leverages the low-rank approximation method to reduce the communication cost, has global superlinear convergence. At each iteration, both DAN and DAN-LA need $n - 1$ communication rounds, and the communication cost for each communication round is $O(d^2)$ and $O(d)$ respectively.

We compare the convergence rate and per-iteration communication cost with these Newton-type methods in Table 2. Note that the first five methods in the table have rates that depend on the condition number of the problem, and as such, do not have the benefits normally attributed to pure Newton’s method. Note also that the two prior methods which do have rates independent of the condition number have high cost of communication. Our methods combine the best of both worlds, and are the only methods we know about which do so: we obtain fast rates independent of the condition number, and allow for $O(d)$ communication per communication round. We were able to achieve this by a complete redesign of how second order methods should work in the distributed setting. Our methods are not simple extensions of existing schemes, and our proofs use novel arguments and techniques.

2 Three Steps Towards an Efficient Distributed Newton Type Method

In order to better explain the algorithms and results of this paper, we will proceed through several steps in a gradual explanation of the ideas that ultimately lead to our methods. While this is not the process we used to come up with our methods, in retrospect we believe that our methods and results will be understood more easily when seen as having been arrived at in this way. In other words, we have constructed what we believe is a plausible discovery story, one enabling faster and better comprehension. If these ideas seem to follow naturally, it is because we made a conscious effort to make them appear that way. The goal of this paper is to develop communication efficient variants of Newton’s method for solving the distributed optimization problem (1).

2.1 Naive distributed implementation of Newton’s method

Newton’s method applied to problem (1) performs the iteration

$$x^{k+1} = x^k - (\nabla^2 P(x^k))^{-1} \nabla P(x^k) \stackrel{(1)}{=} x^k - (\mathbf{H}(x^k) + \lambda \mathbf{I})^{-1} \nabla P(x^k). \quad (9)$$

A naive way to implement this method in the *parameter server* framework is for each node i to compute the Hessian $\mathbf{H}_i(x^k)$ and gradient $\nabla f_i(x^k)$ and to communicate these objects to the server. The server then averages the local Hessians $\mathbf{H}_i(x^k)$ to produce $\mathbf{H}(x^k)$ via (8), and averages the local gradients $\nabla f_i(x^k)$ to produce $\nabla f(x^k)$. The server then adds

$\lambda \mathbf{I}$ to the Hessian, producing $\mathbf{H}(x^k) + \lambda \mathbf{I} = \nabla^2 P(x^k)$, adds λx^k to the gradient, producing $\nabla P(x^k) = \nabla f(x^k) + \lambda x^k$, and subsequently performs the Newton step (9). The resulting vector x^{k+1} is then broadcasted to the nodes and the process is repeated.

This implementation mirrors the way GD and many other first order methods are implemented in the parameter server framework. However, unlike in the case of GD, where only $O(d)$ floats need to be sent and received by each node in each iteration, the upstream communication in Newton’s method requires $O(d^2)$ floats to be communicated by each worker to the server. Since d is typically very large, this is prohibitive in practice. Moreover, computation of the Newton’s step by the parameter server is much more expensive than simple averaging of the gradients performed by gradient type methods. However, in this paper we will not be concerned with the cost of the Newton step itself, as we will assume the server is powerful enough and the network connection is slow enough for this step not to be the main bottleneck of the iteration. Instead, we assume that the communication steps in general, and the $O(d^2)$ communication of the Hessian matrices in particular, is what forms the bottleneck. The $O(d)$ per node communication cost of the local gradients is negligible, and so is the $O(d)$ broadcast of the updated model.

2.2 A better implementation taking advantage of the structure of $\mathbf{H}_{ij}(x)$

The above naive implementation can be improved in the setting when $m < d^2$ by taking advantage of the explicit structure (7) of the local Hessians as a conic combination of positive semidefinite rank one matrices:

$$\mathbf{H}_i(x) = \frac{1}{m} \sum_{j=1}^m h_{ij}(x) a_{ij} a_{ij}^\top. \quad (10)$$

Indeed, assuming that the server has direct access to all the training data vectors $a_{ij} \in \mathbb{R}^d$ (these vectors can be sent to the server at the start of the process), node i can send the m coefficients $h_{i1}(x), \dots, h_{im}(x)$ to the server instead, and the server is then able to reconstruct the Hessian matrix $\mathbf{H}_i(x)$ from this information. This way, each node sends $O(m + d)$ floats to the server, which is a substantial improvement on the naive implementation in the regime when $m \ll d^2$. However, when $m \gg d$, the upstream communication cost is still substantially larger than the $O(d)$ cost of GD. If the server does not have enough memory to store all vectors a_{ij} , this procedure does not work.

2.3 NEWTON-STAR: Newton’s method with a single Hessian

We now introduce a simple idea which, surprisingly, enables us to *remove the need to iteratively communicate any coefficients altogether*. Assume, for the sake of argument, that we know the values $h_{ij}(x^*)$ for all i, j . That is, assume the server has access to coefficients $h_{ij}(x^*)$ for all i, j , and that each node i has access to coefficients $h_{ij}(x^*)$ for $j = 1, \dots, m$, i.e., to the vector

$$h_i(x) := (h_{i1}(x), \dots, h_{im}(x)) \in \mathbb{R}^m \quad (11)$$

for $x = x^*$. Next, consider the following new Newton-like method which we call **NEWTON-STAR (NS)**, where the “star” points to the method’s reliance on the knowledge of the

optimal solution x^* :

$$\begin{aligned} x^{k+1} &= x^k - (\nabla^2 P(x^*))^{-1} \nabla P(x^k) \\ &\stackrel{(1)}{=} x^k - (\mathbf{H}(x^*) + \lambda \mathbf{I})^{-1} \left(\frac{1}{n} \sum_{i=1}^n \nabla f_i(x^k) + \lambda x^k \right). \end{aligned} \quad (12)$$

Since the server knows $\mathbf{H}(x^*)$, all that the nodes need to communicate are the local gradients $\nabla f_i(x^k)$, which costs $O(d)$ per node. The server then computes x^{k+1} , broadcasts it back to the nodes, and the process is repeated. This method has the same per-iteration $O(d)$ communication complexity as GD. However, as we show next, the number of iterations (which is the same as the number of communications) of **NEWTON-STAR** does not depend on the condition number – a property it borrows from the classical Newton’s method. The following theorem says that **NEWTON-STAR** enjoys *local quadratic convergence*.

Theorem 2.1 (Local quadratic convergence). *Let Assumption 1.1 hold, and assume that $\mathbf{H}(x^*) \succeq \mu^* \mathbf{I}$ for some $\mu^* \geq 0$ (for instance, this holds if f is μ^* -strongly convex) and that $\mu^* + \lambda > 0$. Then for any starting point $x^0 \in \mathbb{R}^d$, the iterates of **NEWTON-STAR** for solving problem (1) satisfy the following inequality:*

$$\|x^{k+1} - x^*\| \leq \frac{\nu}{2(\mu^* + \lambda)} \cdot \left(\frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m \|a_{ij}\|^3 \right) \cdot \|x^k - x^*\|^2. \quad (13)$$

Proof. By the first order optimality conditions, we have

$$\nabla f(x^*) + \lambda x^* = 0. \quad (14)$$

Let $\mathbf{H}_* := \mathbf{H}(x^*)$. Since $\mathbf{H}_* \succeq \mu^* \mathbf{I}$, we have $\mathbf{H}_* + \lambda \mathbf{I} \succeq (\mu^* + \lambda) \mathbf{I}$, and hence

$$\|(\mathbf{H}_* + \lambda \mathbf{I})^{-1}\| \leq \frac{1}{\mu^* + \lambda}. \quad (15)$$

Using (14) and (15) and subsequently applying Jensen’s inequality to the function $x \mapsto \|x\|$, we get

$$\begin{aligned} \|x^{k+1} - x^*\| &= \left\| x^k - x^* - (\mathbf{H}_* + \lambda \mathbf{I})^{-1} \nabla P(x^k) \right\| \\ &\stackrel{(14)}{=} \left\| (\mathbf{H}_* + \lambda \mathbf{I})^{-1} [(\mathbf{H}_* + \lambda \mathbf{I})(x^k - x^*) - (\nabla f(x^k) - \nabla f(x^*) + \lambda(x^k - x^*))] \right\| \\ &\stackrel{(15)}{\leq} \frac{1}{\mu^* + \lambda} \left\| (\mathbf{H}_* + \lambda \mathbf{I})(x^k - x^*) - (\nabla f(x^k) - \nabla f(x^*)) - \lambda(x^k - x^*) \right\| \\ &= \frac{1}{\mu^* + \lambda} \left\| \frac{1}{n} \sum_{i=1}^n \mathbf{H}_i(x^*)(x^k - x^*) - \frac{1}{n} \sum_{i=1}^n (\nabla f_i(x^k) - \nabla f_i(x^*)) \right\| \\ &\leq \frac{1}{n(\mu^* + \lambda)} \sum_{i=1}^n \left\| \mathbf{H}_i(x^*)(x^k - x^*) - (\nabla f_i(x^k) - \nabla f_i(x^*)) \right\| \\ &\stackrel{(7)}{=} \frac{1}{n(\mu^* + \lambda)} \sum_{i=1}^n \left\| \frac{1}{m} \sum_{j=1}^m [h_{ij}(x^*) a_{ij} a_{ij}^\top (x^k - x^*) - (\nabla f_{ij}(x^k) - \nabla f_{ij}(x^*))] \right\|. \end{aligned} \quad (16)$$

We now use the fundamental theorem of calculus to express difference of gradients $\nabla f_{ij}(x^k) - \nabla f_{ij}(x^*)$ in an integral, obtaining

$$\nabla f_{ij}(x^k) - \nabla f_{ij}(x^*) = \int_0^1 \nabla^2 f_{ij}(x^* + \tau(x^k - x^*))(x^k - x^*) d\tau. \quad (17)$$

Plugging this representation into (16) and noting that $\nabla^2 f_{ij}(x) \equiv \mathbf{H}_{ij}(x)$ (see (4)), we can continue:

$$\begin{aligned}
\|x^{k+1} - x^*\| &\stackrel{(16)+(17)}{\leq} \frac{1}{n(\mu^* + \lambda)} \sum_{i=1}^n \left\| \frac{1}{m} \sum_{j=1}^m (h_{ij}(x^*) a_{ij} a_{ij}^\top (x^k - x^*) \right. \\
&\quad \left. - \int_0^1 \mathbf{H}_{ij}(x^* + \tau(x^k - x^*)) (x^k - x^*) d\tau \right\| \\
&\stackrel{(4)}{=} \frac{1}{n(\mu^* + \lambda)} \sum_{i=1}^n \left\| \frac{1}{m} \sum_{j=1}^m (h_{ij}(x^*) a_{ij} a_{ij}^\top (x^k - x^*) \right. \\
&\quad \left. - \int_0^1 h_{ij}(x^* + \tau(x^k - x^*)) a_{ij} a_{ij}^\top (x^k - x^*) d\tau \right\| \\
&= \frac{1}{n(\mu^* + \lambda)} \sum_{i=1}^n \left\| \frac{1}{m} \sum_{j=1}^m a_{ij} a_{ij}^\top (x^k - x^*) \left(h_{ij}(x^*) - \int_0^1 h_{ij}(x^* + \tau(x^k - x^*)) d\tau \right) \right\| \\
&\leq \frac{\|x^k - x^*\|}{(\mu^* + \lambda)} \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m \|a_{ij}\|^2 \left| \int_0^1 h_{ij}(x^*) - h_{ij}(x^* + \tau(x^k - x^*)) d\tau \right|. \tag{18}
\end{aligned}$$

In the last step we have again used Jensen's inequality applied to the function $x \mapsto \|x\|$, followed by inequalities of the form $\|\mathbf{A}_{ij} x t_{ij}\| \leq \|\mathbf{A}_{ij}\| \|x\| |t_{ij}|$ for $\mathbf{A}_{ij} = a_{ij} a_{ij}^\top$, $x = x^k - x^*$ and $t_{ij} \in \mathbb{R}$.

From (6) we obtain $|h_{ij}(x^*) - h_{ij}(x^* + \tau(x^k - x^*))| \leq \nu \tau \|a_{ij}\| \cdot \|x^k - x^*\|$, which implies that

$$\left| \int_0^1 h_{ij}(x^*) - h_{ij}(x^* + \tau(x^k - x^*)) d\tau \right| \leq \int_0^1 \nu \tau \|a_{ij}\| \cdot \|x^k - x^*\| d\tau = \frac{\nu \|a_{ij}\|}{2} \cdot \|x^k - x^*\|.$$

Plugging this into (18), we finally arrive at (13). \square

Note that we do not need to assume f to be convex or strongly convex. All we need to assume is positive definiteness of the Hessian at the optimum. This implies local strong convexity, and since our convergence result is local, that is all we need.

Remark. Besides NEWTON-STAR, we have designed another new Newton-type method with a local quadratic rate. This method, which we call MAX-NEWTON, is similar to NEWTON-STAR in that it relies on the knowledge of the coefficients $h_{ij}(x^*)$ for $j = 1, \dots, m$. We describe this method in Appendix D.

3 NEWTON-LEARN: Learning the Hessian and Local Convergence Theory

In Sections 2.1, 2.2 and 2.3 we have gone through three steps in our story, with the first true innovation and contribution of this paper being the NEWTON-STAR method and its rate. We have now sufficiently prepared the ground to motivate our first *key* contribution: the NEWTON-LEARN method. We only outline the basic insights behind this method here; the details are included in Section 3.

3.1 The main iteration

In NEWTON-LEARN we maintain a sequence of vectors

$$h_i^k = (h_{i1}^k, \dots, h_{im}^k) \in \mathbb{R}^m,$$

for all $i = 1, \dots, n$ throughout the iterations $k \geq 0$ with the goal of *learning* the values $h_{ij}(x^*)$ for all i, j . That is, we construct the sequence with the explicit intention to enforce

$$h_{ij}^k \rightarrow h_{ij}(x^*) \quad \text{as} \quad k \rightarrow \infty. \quad (19)$$

Using $h_{ij}^k \approx h_{ij}(x^*)$ we estimate the Hessian $\mathbf{H}(x^*)$ via

$$\mathbf{H}(x^*) \approx \mathbf{H}^k := \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m h_{ij}^k a_{ij} a_{ij}^\top, \quad (20)$$

and then perform a similar iteration to (12):

$$x^{k+1} = x^k - (\mathbf{H}^k + \lambda \mathbf{I})^{-1} \nabla P(x^k). \quad (21)$$

3.2 Learning the coefficients: the idea

To complete the description of the method, we need to explain how the vectors h_i^{k+1} are updated. This is also the place where we can force the method to be communication efficient. Indeed, if we can design a rule that would enforce the update vectors $h_i^{k+1} - h_i^k$ to be *sparse*, say

$$\|h_i^{k+1} - h_i^k\|_0 \leq s \quad (22)$$

for some $1 \leq s \leq m$ and all i and k , then the upstream communication by each node in each iteration would be of the order $O(s + d)$ only (provided the server has access to all vectors a_{ij})! That is, each node i only needs to communicate s entries of the update vector $h_i^{k+1} - h_i^k$ as the rest are equal to zero, and each node also needs to communicate the d dimensional gradient $\nabla f_i(x^k)$. Note that $O(s + d)$ can be interpreted as an *interpolation* of the $O(m + d)$ per-iteration communication complexity of the structure-aware implementation of Newton's method from Section 2.2, and of the $O(d)$ per-iteration communication complexity of NEWTON-STAR described in Section 2.3.

In the more realistic regime when the server does *not* have access to the data $\{a_{ij}\}$, we ask each worker i to additionally send the corresponding s vectors a_{ij} , which costs extra $O(sd)$ in communication per node. However, when $s = O(1)$, this is the same per-iteration communication effort as that of GD.

We develop two different update rules defining the evolution of the vectors h_1^k, \dots, h_n^k . This first rule (see (25)) applies to the $\lambda > 0$ case and leads to our first variant of NEWTON-LEARN which we call NL1 (see Algorithm 1). This rule and the method are described in Section 3.5. The second rule applies also to the $\lambda = 0$ cases and leads to our second variant of NEWTON-LEARN which we call NL2 (see Algorithm 2). This rule and the method are described in Section 3.6.

3.3 Outline of fast local convergence theory

We show in Theorem 3.2 (covering NL1) and Theorem 3.5 (covering NL2) that NEWTON-LEARN enjoys a local linear rate wrt a certain Lyapunov function which involves the term $\|x^k - x^*\|^2$ and also all terms of the form $\|h_i^k - h_i(x^*)\|^2$. This means that i) the main iteration (21) works, i.e., x^k converges to x^* at a local linear rate, and that ii) the learning procedure works, and the desired convergence described in (19) occurs at a local linear rate.

In addition, we also establish a local superlinear rate of $\|x^k - x^*\|^2$. Remarkably, these rates are *independent of any condition number, which is in sharp contrast with virtually all results on distributed Newton-type methods we are aware of*.

Moreover, we wish to remark that second order methods are not typically analyzed using a Lyapunov style analysis. Indeed, we only know of a couple works that do so. First, Kovalev et al. [21] develop stochastic Newton and cubic Newton methods of a different structure and scope from ours. They do not consider distributed optimization nor communication compression. Second, Kovalev et al. [22] develop a stochastic BFGS method. Again, their method and scope is very different from ours. Hence, *our analysis may be of independent interest as it adds to the arsenal of theoretical tools which could be used in a more precise analysis of other second order methods*.

3.4 Compressed learning

Instead of merely relying on sparse updates for the vectors h_i^k (see (22)), we provide a more general communication compression strategy which includes sparsification as a special case [1]. We do so via the use of a *random compression* operator. We say that a randomized map $\mathcal{C} : \mathbb{R}^m \rightarrow \mathbb{R}^m$ is a *compression operator (compressor)* if there exists a constant $\omega \geq 0$ such that the following relations hold for all $x \in \mathbb{R}^m$:

$$\mathbb{E}[\mathcal{C}(x)] = x \quad (23)$$

$$\mathbb{E}\|\mathcal{C}(x)\|^2 \leq (\omega + 1)\|x\|^2. \quad (24)$$

The identity compressor $\mathcal{C}(x) \equiv x$ satisfies these relations with $\omega = 0$. The larger the variance parameter ω is allowed to be, the easier it can be to construct a compressor \mathcal{C} for which the value $\mathcal{C}(x)$ can be encoded using a small number of bits only. We refer the reader to [5] for a list of several compressors and their properties.

3.5 NL1 (learning in the $\lambda > 0$ case)

We now consider the case where all loss functions φ_{ij} are convex and $\lambda > 0$.

Assumption 3.1. *Each φ_{ij} is convex, $\lambda > 0$.*

When combined with Assumption 1.1, Assumption 3.1 implies that $\varphi_{ij}''(t) \geq 0$ for all t , hence $h_{ij}(x) = \varphi_{ij}''(a_i^\top x) \geq 0$ for all $x \in \mathbb{R}^d$. In particular, $h_{ij}(x^*) \geq 0$ for all i, j . Since we wish to construct a sequence of vectors $h_i^k = (h_{i1}^k, \dots, h_{im}^k) \in \mathbb{R}^m$ satisfying $h_{ij}^k \rightarrow h_{ij}(x^*)$, it makes sense to try to enforce all vectors in this sequence to have *nonnegative entries*:

$$h_{ij}^k \geq 0.$$

Since \mathbf{H}^k arises as a linear combination of the rank-one matrices $a_{ij}a_{ij}^\top$ (see (20)), this makes \mathbf{H}^k positive semidefinite, which in turn means that the matrix $\mathbf{H}^k + \lambda\mathbf{I}$ appearing in the main iteration (21) of NEWTON-LEARN is invertible, and hence the iteration is *well defined*.⁵

⁵Positive definiteness of Hessian estimates is enforced in several popular quasi-Newton methods as well; for instance, in the BFGS method [6, 9, 11, 43]. However, quasi-Newton methods operate in a markedly different manner, and the way in which positive definiteness is enforced there is also different.

3.5.1 The learning iteration and the NL1 algorithm

In particular, in **NEWTON-LEARN** each node i computes the vector $h_i(x^k) \in \mathbb{R}^m$ of second derivatives defined in (11), and then performs the update

$$\boxed{h_i^{k+1} = [h_i^k + \eta \mathcal{C}_i^k(h_i(x^k) - h_i^k)]_+,} \quad (25)$$

where $\eta > 0$ is a learning rate, \mathcal{C}_i^k is a freshly sampled compressor by node i at iteration k . By $[\cdot]_+$ we denote the positive part function applied element-wise, defined for scalars as follows: $[t]_+ = t$ if $t \geq 0$ and $[t]_+ = 0$ otherwise.

We remark that it is possible to interpret the learning procedure (25) as one step of projected stochastic gradient descent (SGD) applied to a certain quadratic optimization problem whose unique solution is the vector $h_i(x^k)$.

The **NL1** algorithm (Algorithm 1) arises as the combination of the Newton-like update (21) (adjusted to take account of the explicit regularizer) and the learning procedure (25). It is easy to see that the update rule for \mathbf{H}^k in **NL1** is designed to ensure that \mathbf{H}^k remains of the form $\mathbf{H}^k = \frac{1}{n} \sum_{i=1}^n \mathbf{H}_i^k$, where $\mathbf{H}_i^k = \frac{1}{m} \sum_{j=1}^m h_{ij}^k a_{ij} a_{ij}^\top$. The update rule for x^k , performed by the server, is identical to (21), with an extra provision for the regularizer. The vector x^{k+1} is broadcasted to all workers. Let us comment on how the key communication step is implemented. If the server does not have direct access to the training data vectors $\{a_{ij}\}$, we choose Option 1, otherwise we choose Option 2. A key property of **NL1** is that the server is able to maintain copies of the learning vectors h_i^k *without the need for these vectors to be communicated by the workers to the server*. Indeed, provided the workers and the server agree on the same set of initial vectors h_1^0, \dots, h_n^0 , update (25) can be independently computed by the server as well from its memory state h_i^k and the compressed message $\mathcal{C}_i^k(h_i(x^k) - h_i^k)$ received from node i . This strategy is reminiscent of the way the key step in the first-order method DIANA [16, 29] is executed. In this sense, **NL1** can be seen as arising from a successful marriage of Newton's method and the DIANA trick.

3.5.2 Theory

In our theoretical results we rely on the Lyapunov function

$$\Phi_1^k := \|x^k - x^*\|^2 + \frac{1}{3mn\eta\nu^2 R^2} \mathcal{H}^k, \quad \mathcal{H}^k := \sum_{i=1}^n \|h_i^k - h_i(x^*)\|^2.$$

Our main theorem follows.

Theorem 3.2 (Convergence of **NL1**). *Let Assumptions 1.1 and 3.1 hold. Let $\eta \leq \frac{1}{\omega+1}$ and assume that $\|x^k - x^*\|^2 \leq \frac{\lambda^2}{12\nu^2 R^6}$ for all $k \geq 0$. Then for Algorithm 1 we have the inequalities*

$$\begin{aligned} \mathbb{E}[\Phi_1^k] &\leq \theta_1^k \Phi_1^0, \\ \mathbb{E} \left[\frac{\|x^{k+1} - x^*\|^2}{\|x^k - x^*\|^2} \right] &\leq \theta_1^k \left(6\eta + \frac{1}{2} \right) \frac{\nu^2 R^6}{\lambda^2} \Phi_1^0, \end{aligned}$$

where $\theta_1 := 1 - \min \left\{ \frac{\eta}{2}, \frac{5}{8} \right\}$.

Algorithm 1 NL1: NEWTON-LEARN ($\lambda > 0$ case)

Parameters: learning rate $\eta > 0$

Initialization: $x^0 \in \mathbb{R}^d$; $h_1^0, \dots, h_n^0 \in \mathbb{R}_+^m$; $\mathbf{H}^0 = \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m h_{ij}^0 a_{ij} a_{ij}^\top \in \mathbb{R}^{d \times d}$

for $k = 0, 1, 2, \dots$ **do**

 Broadcast x^k to all workers

for each node $i = 1, \dots, n$ **do**

 Compute local gradient $\nabla f_i(x^k)$

$h_i^{k+1} = [h_i^k + \eta \mathcal{C}_i^k(h_i(x^k) - h_i^k)]_+$

 Send $\nabla f_i(x^k)$ and $\mathcal{C}_i^k(h_i(x^k) - h_i^k)$ to server

Option 1: Send $\{a_{ij} : h_{ij}^{k+1} - h_{ij}^k \neq 0\}$ to server

Option 2: Do nothing if server knows $\{a_{ij} : \forall j\}$

end for

$x^{k+1} = x^k - (\mathbf{H}^k + \lambda \mathbf{I})^{-1} \left(\frac{1}{n} \sum_{i=1}^n \nabla f_i(x^k) + \lambda x^k \right)$

$\mathbf{H}^{k+1} = \mathbf{H}^k + \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m (h_{ij}^{k+1} - h_{ij}^k) a_{ij} a_{ij}^\top$

end for

Since the stepsize bound $\eta \leq \frac{1}{\omega+1}$ is independent of the condition number, the linear convergence rates of $\mathbb{E}[\Phi_1^k]$ and $\mathbb{E} \left[\frac{\|x^{k+1} - x^*\|^2}{\|x^k - x^*\|^2} \right]$ are both *independent of the condition number*. Next, we explore under what conditions we can guarantee for all the iterates to stay in a small neighborhood.

Lemma 3.3. *Let Assumptions 1.1 and 3.1 hold. Assume h_{ij}^k is a convex combination of $\{h_{ij}(x^0), h_{ij}(x^1), \dots, h_{ij}(x^k)\}$ for all i, j and k . Assume $\|x^0 - x^*\|^2 \leq \frac{\lambda^2}{12\nu^2 R^6}$. Then*

$$\|x^k - x^*\|^2 \leq \frac{\lambda^2}{12\nu^2 R^6} \quad \text{for all } k \geq 0.$$

It is easy to verify that if we choose $h_{ij}^0 = h_{ij}(x^0)$ and use the random sparsification compressor and $\eta \leq \frac{1}{\omega+1}$, then h_{ij}^k is always a convex combination of $\{h_{ij}(x^0), h_{ij}(x^1), \dots, h_{ij}(x^k)\}$ for $k \geq 0$. Thus, from Lemma 3.3 we can guarantee that all the iterates stay in the small neighborhood assumed in Theorem 3.2 as long as the initial point x^0 is in it.

3.6 NL2 (learning in the $\lambda \geq 0$ case)

In this subsection, we consider the case where P is μ -strongly convex. Note that we do not require the components f_{ij} to be convex.

Assumption 3.4. *P is μ -strongly convex, $|h_{ij}^k| \leq \gamma$ for $k \geq 0$.*

3.6.1 The learning iteration and the NL2 algorithm

As in Algorithm 1, we use a sequence of vectors $\{h_i^k\}_{k \geq 0}$ to learn $h_i(x^*)$. However, this time we rely on a different technique for enforcing positive definiteness of the Hessian estimator. Since λ can be zero, our previous technique aimed at forcing the coefficients

h_{ij}^k to be nonnegative will not work. So, we give up on this, and instead of (25) we use the simpler update

$$\boxed{h_i^{k+1} = h_i^k + \eta \mathcal{C}_i^k(h_i(x^k) - h_i^k).} \quad (26)$$

In order to guarantee positive definiteness of the Hessian estimator $\mathbf{H}^k + \lambda \mathbf{I}$ we instead rely on the second part of Assumption 3.4. Provided that there exists $\gamma > 0$ such that $|h_{ij}^k| \leq \gamma$ for all i, j , note that $\frac{h_{ij}(x^k) + 2\gamma}{h_{ij}^k + 2\gamma}$ is always positive. Noticing that each $a_{ij}a_{ij}^\top$ is positive semidefinite and that $\nabla^2 f(x^k)$ can be expressed in the form

$$\frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m \left(\frac{h_{ij}(x^k) + 2\gamma}{h_{ij}^k + 2\gamma} \cdot (h_{ij}^k + 2\gamma) - 2\gamma \right) a_{ij}a_{ij}^\top,$$

for $\beta^k := \max_{i,j} \frac{h_{ij}(x^k) + 2\gamma}{h_{ij}^k + 2\gamma}$, we get the inequality

$$\begin{aligned} & \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m [\beta^k (h_{ij}^k + 2\gamma) - 2\gamma] a_{ij}a_{ij}^\top - \nabla^2 f(x^k) \\ &= \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m \left[\beta^k - \frac{h_{ij}(x^k) + 2\gamma}{h_{ij}^k + 2\gamma} \right] (h_{ij}^k + 2\gamma) a_{ij}a_{ij}^\top \succeq \mathbf{0}, \end{aligned}$$

where $\mathbf{0}$ is the $d \times d$ zero matrix, and $\mathbf{A} \succeq \mathbf{B}$ means $\mathbf{A} - \mathbf{B}$ is positive semidefinite. Thus, if we can maintain the Hessian estimator in the form

$$\mathbf{H}^k := \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m [\beta^k (h_{ij}^k + 2\gamma) - 2\gamma] a_{ij}a_{ij}^\top,$$

then $\mathbf{H}^k + \lambda \mathbf{I} \succeq \nabla^2 f(x^k) + \lambda \mathbf{I} = \nabla^2 P(x^k) \succeq \mu \mathbf{I}$, where the last inequality follows from Assumption 3.4. To achieve this goal, we use an auxiliary matrix \mathbf{A}^k , and maintain $\mathbf{A}^k = \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m (h_{ij}^k + 2\gamma) a_{ij}a_{ij}^\top$, and $\mathbf{H}^k = \beta^k \mathbf{A}^k - 2\gamma \cdot \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m a_{ij}a_{ij}^\top$. The rest of Algorithm 2 is the same as Algorithm 1.

3.6.2 Theory

Our analysis of NL2 relies on the Lyapunov function

$$\Phi_2^k := \|x^k - x^*\|^2 + \frac{1}{3mn\eta\nu^2 R^2} \mathcal{H}^k, \quad \mathcal{H}^k := \sum_{i=1}^n \|h_i^k - h_i(x^*)\|^2.$$

We now present our main convergence result for NL2.

Theorem 3.5 (Convergence of NL2). *Let Assumptions 1.1 and 3.4 hold. Assume $\eta \leq \frac{1}{\omega+1}$ and $\|x^k - x^*\|^2 \leq \frac{\mu^2}{432mn\nu^2 R^6}$ for all $k \geq 0$. Then for Algorithm 2 we have the inequalities*

$$\begin{aligned} \mathbb{E}[\Phi_2^k] &\leq \theta_2^k \Phi_2^0, \\ \mathbb{E} \left[\frac{\|x^{k+1} - x^*\|^2}{\|x^k - x^*\|^2} \right] &\leq \theta_2^k (3mn\eta + 1) \frac{72\nu^2 R^6}{\mu^2} \Phi_2^0, \end{aligned}$$

where $\theta_2 := 1 - \min \left\{ \frac{\eta}{2}, \frac{1}{2} \right\}$.

Algorithm 2 NL2: NEWTON-LEARN (general case)

Parameters: $\eta > 0; \gamma > 0$

Initialization: $x^0 \in \mathbb{R}^d; h_i^0 \in \mathbb{R}^m; \mathbf{A}^0 = \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m (h_{ij}^0 + 2\gamma) a_{ij} a_{ij}^\top \in \mathbb{R}^{d \times d}$

for $k = 0, 1, 2, \dots$ **do**

 broadcast x^k to all workers

for $i = 1, \dots, n$ **do**

 Compute local gradient $\nabla f_i(x^k)$

$h_i^{k+1} = h_i^k + \eta \mathcal{C}_i^k(h_i(x^k) - h_i^k)$

$\beta_i^k = \max_{j \in [m]} \frac{h_{ij}(x^k) + 2\gamma}{h_{ij}^k + 2\gamma}$

 Send $\nabla f_i(x^k)$, β_i^k , and $\eta \mathcal{C}_i^k(h_i(x^k) - h_i^k)$ to server

Option 1: Send $\{a_{ij} : h_{ij}^{k+1} - h_{ij}^k \neq 0\}$ to server

Option 2: Do nothing if server knows $\{a_{ij} : \forall j\}$

end for

$\beta^k = \max_i \{\beta_i^k\}$

$\mathbf{H}^k = \beta^k \mathbf{A}^k - 2\gamma \cdot \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m a_{ij} a_{ij}^\top \in \mathbb{R}^{d \times d}$

$x^{k+1} = x^k - (\mathbf{H}^k + \lambda \mathbf{I})^{-1} \left(\frac{1}{n} \sum_{i=1}^n \nabla f_i(x^k) + \lambda x^k \right)$

$\mathbf{A}^{k+1} = \mathbf{A}^k + \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m (\eta \mathcal{C}_i^k(h_i(x^k) - h_i^k))_j a_{ij} a_{ij}^\top$

end for

As before, we give sufficient conditions guaranteeing that the iterates stay in a small neighborhood of the optimum.

Lemma 3.6. *Let Assumptions 1.1 and 3.4 hold. Assume h_{ij}^k is a convex combination of $\{h_{ij}(x^0), h_{ij}(x^1), \dots, h_{ij}(x^k)\}$ for all i, j and k . Assume $\|x^0 - x^*\|^2 \leq \frac{\mu^2}{432mn\nu^2 R^6}$. Then*

$$\|x^k - x^*\|^2 \leq \frac{\mu^2}{432mn\nu^2 R^6} \quad \text{for all } k \geq 0.$$

If we choose $h_{ij}^0 = h_{ij}(x^0)$, use a random compressor with variance ω , and choose stepsize $\eta \leq \frac{1}{\omega+1}$, then h_{ij}^k is a convex combination of $\{h_{ij}(x^0), h_{ij}(x^1), \dots, h_{ij}(x^k)\}$ for all $k \geq 0$. Thus, via Lemma 3.6 we can guarantee all the iterates to be in the small neighborhood required by Theorem 3.5 as long as the initial point x^0 is in it.

4 CUBIC-NEWTON-LEARN: Global Convergence Theory via Cubic Regularization

In this section, we develop the CUBIC-NEWTON-LEARN method which can obtain *global* convergence and superlinear convergence. We make the following assumption throughout this section.

Assumption 4.1. *Assume $|h_{ij}^k| \leq \gamma$ for $k \geq 0$, and $\|a_{ij}\| \leq R$ for all i, j . Assume $R_x := \sup_{x \in \mathbb{R}^d} \{\|x - x^*\| : P(x) \leq P(x^0)\} < +\infty$.*

Recall that $\mathbf{H}(x)$ is the Hessian of f at point x . For any $x, y \in \mathbb{R}^d$, we have

$$\begin{aligned}
\|\mathbf{H}(x) - \mathbf{H}(y)\| &= \left\| \frac{1}{n} \sum_{i=1}^n \frac{1}{m} \sum_{j=1}^m (h_{ij}(x) - h_{ij}(y)) a_{ij} a_{ij}^\top \right\| \\
&\leq \frac{1}{n} \sum_{i=1}^n \frac{1}{m} \sum_{j=1}^m \|a_{ij}\|^2 |h_{ij}(x) - h_{ij}(y)| \\
&\stackrel{(6)}{\leq} \frac{1}{n} \sum_{i=1}^n \frac{1}{m} \sum_{j=1}^m \nu \|a_{ij}\|^3 \|x - y\| \leq \nu R^3 \|x - y\|.
\end{aligned}$$

Let $M := \nu R^3$. Then f has M -Lipschitz continuous Hessian, and so does P .

4.1 CNL: the algorithm

At a high level, our CNL method (Algorithm 3) can be seen as the “combination” of Algorithm 2 (NL2) and the cubic Newton method of Griewank [13], Nesterov and Polyak [33]. We use the same learning procedure for h_i^k , and the same construction for obtaining a Hessian estimator \mathbf{H}^k for f as those used in Algorithm 2. In particular, since $|h_{ij}^k| \leq \gamma$, same as Algorithm 2, we have

$$\mathbf{A}^k = \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m (h_{ij}^k + 2\gamma) a_{ij} a_{ij}^\top$$

and

$$\mathbf{H}^k = \max_{ij} \frac{h_{ij}(x^k) + 2\gamma}{h_{ij}^k + 2\gamma} \frac{1}{n} \sum_{i=1}^n \frac{1}{m} \sum_{j=1}^m (h_{ij}^k + 2\gamma) a_{ij} a_{ij}^\top - 2\gamma \cdot \frac{1}{n} \sum_{i=1}^n \frac{1}{m} \sum_{j=1}^m a_{ij} a_{ij}^\top \succeq \nabla^2 f(x^k).$$

Thus, $\mathbf{H}^k + \lambda \mathbf{I} \succeq \nabla^2 f(x^k) + \lambda \mathbf{I} = \nabla^2 P(x^k)$. If we let

$$\mathbf{H}_i^k := \beta^k \frac{1}{m} \sum_{j=1}^m (h_{ij}^k + 2\gamma) a_{ij} a_{ij}^\top - 2\gamma \cdot \frac{1}{m} \sum_{j=1}^m a_{ij} a_{ij}^\top,$$

then $\mathbf{H}^k = \frac{1}{n} \sum_{i=1}^n \mathbf{H}_i^k$. Since P has M -Lipschitz continuous Hessian, we have

$$P(x^k + s) \leq P(x^k) + T(x^k, s), \quad \text{for all } x \in \mathbb{R}^d,$$

where

$$T(x^k, s) := \langle \nabla f(x^k) + \lambda x^k, s \rangle + \frac{1}{2} \langle (\mathbf{H}^k + \lambda \mathbf{I}) s, s \rangle + \frac{M}{6} \|s\|^3.$$

Finally, the search direction s^k is obtained by minimizing $T(x^k, s)$. This subproblem can be solved by computing the eigenvalue decomposition and then solving a one-dimensional nonlinear equation [12, 14]. The details can be found in Section C.1.

Algorithm 3 CNL: CUBIC-NEWTON-LEARN

Parameters: $\eta > 0; \gamma > 0$

Initialization: $x^0 \in \mathbb{R}^d; h_i^0 \in \mathbb{R}^m; \mathbf{A}^0 = \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m (h_{ij}^0 + 2\gamma) a_{ij} a_{ij}^\top \in \mathbb{R}^{d \times d}$

for $k = 0, 1, 2, \dots$ **do**

 broadcast x^k to all workers

for $i = 1, \dots, n$ **do**

 Compute local gradient $\nabla f_i(x^k)$

 Update the vector of coefficients $h_i^{k+1} = h_i^k + \eta \mathcal{C}_i^k(h_i(x^k) - h_i^k)$

$\beta_i^k = \max_{j \in [m]} \frac{h_{ij}(x^k) + 2\gamma}{h_{ij}^k + 2\gamma}$

 Send $\nabla f_i(x^k)$, β_i^k , and $\mathcal{C}_i^k(h_i(x^k) - h_i^k)$ to server

Option 1: Send $\{a_{ij} : h_{ij}^{k+1} - h_{ij}^k \neq 0\}$ to server

Option 2: Do nothing if server knows $\{a_{ij} : \forall j\}$

end for

$\beta^k = \max_i \{\beta_i^k\}$

$\mathbf{H}^k = \beta^k \mathbf{A}^k - 2\gamma \cdot \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m a_{ij} a_{ij}^\top \in \mathbb{R}^{d \times d}$

$s^k = \operatorname{argmin}_{s \in \mathbb{R}^d} T(x^k, s)$

$x^{k+1} = x^k + s^k$

$\mathbf{A}^{k+1} = \mathbf{A}^k + \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m (\eta \mathcal{C}_i^k(h_i(x^k) - h_i^k))_j a_{ij} a_{ij}^\top$

end for

4.2 Global convergence

The proof of global convergence of CUBIC-NEWTON-LEARN was inspired by that in the stochastic subspace cubic Newton method of Hanzely et al. [14]. First, we need the following technical lemma, which show the progress for each step of Algorithm 3.

Lemma 4.2. *Let Assumption 1.1 and Assumption 4.1 hold. Then for all $k \geq 0$ and any $y \in \mathbb{R}^d$, we have*

$$P(x^{k+1}) \leq P(y) + 9\gamma R^2 \|y - x^k\|^2 + \frac{M}{3} \|y - x^k\|^3.$$

We first consider the general convex case.

Theorem 4.3 (Global convergence of CNL: convex case). *Let Assumption 1.1 and Assumption 4.1 hold and further assume that P is convex. Then for all $k \geq 0$ we have*

$$P(x^k) - P(x^*) \leq \frac{81\gamma R^2 R_x^2}{k} + \frac{9MR_x^3}{k^2} + \frac{4[P(x^0) - P(x^*)]}{k^3}.$$

In particular,

$$k = O\left(\frac{81\gamma R^2 R_x^2}{\epsilon} + \sqrt{\frac{MR_x^3}{\epsilon}} + \left(\frac{P(x^0) - P(x^*)}{\epsilon}\right)^{1/3}\right) \Rightarrow P(x^k) - P(x^*) \leq \epsilon.$$

In the strongly convex case, we have the following linear convergence result.

Theorem 4.4 (Global convergence of CNL: strongly convex case). *Let Assumption 1.1 and Assumption 4.1 hold and assume that P is μ -strongly convex. Then*

$$k = O\left(\left(\frac{\gamma R^2}{\mu} + \sqrt{\frac{MR_x}{\mu}} + 1\right) \log\left(\frac{P(x^0) - P(x^*)}{\epsilon}\right)\right) \Rightarrow P(x^k) - P(x^*) \leq \epsilon.$$

4.3 Superlinear convergence

Let us define the Lyapunov function

$$\Phi_3^k := \|x^k - x^*\|^2 + \frac{4}{9mn\eta\nu^2 R^2} \mathcal{H}^k, \quad \mathcal{H}^k := \sum_{i=1}^n \|h_i^k - h_i(x^*)\|^2.$$

We now show that this Lyapunov function converges locally linearly, and that the squared distance to the optimum converges locally superlinearly. Both rates are independent of any condition number. Note that this also means that our learning procedure works, i.e., the estimates h_i^k converge to the values $h_i(x^*)$, which also means that $\mathbf{H}^k \rightarrow \mathbf{H}(x^*)$.

Theorem 4.5 (Local superlinear convergence of CNL: strongly convex case). *Let Assumption 1.1 and Assumption 4.1 hold. Assume P is μ -strongly convex, and that $\|x^k - x^*\|^2 \leq \frac{\mu^2}{432mn\nu^2 R^6}$ for all $k \geq 0$. Then for Algorithm 3, we have the relations*

$$\begin{aligned} \mathbb{E}[\Phi_3^k] &\leq \theta_3^k \Phi_3^0 \\ \mathbb{E}\left[\frac{\|x^{k+1} - x^*\|^2}{\|x^k - x^*\|^2}\right] &\leq \theta_3^k (3mn\eta + 1) \frac{100\nu^2 R^6}{\mu^2} \Phi_3^0, \end{aligned}$$

where $\theta_3 := 1 - \min\left\{\frac{\eta}{2}, \frac{1}{4}\right\}$.

Theorem 4.4, Theorem 4.5, and the fact that $\frac{\mu}{2}\|x - x^*\|^2 \leq P(x) - P(x^*)$ when P is μ -strongly convex, imply the following corollary.

Corollary 4.6. *Let Assumption 1.1 and Assumption 4.1 hold. Assume P is μ -strongly convex. Then*

$$\lim_{k \rightarrow +\infty} \mathbb{E}\left[\frac{\|x^{k+1} - x^*\|^2}{\|x^k - x^*\|^2}\right] = 0.$$

5 Experiments

We now study the empirical performance of our second order methods NL1, NL2 and CNL, and compare them with relevant benchmarks and with state-of-the-art methods. We test on the regularized logistic regression problem

$$\min_{x \in \mathbb{R}^d} \left\{ \frac{1}{n} \sum_{i=1}^n \frac{1}{m} \sum_{j=1}^m \log(1 + \exp(-b_{ij} a_{ij}^\top x)) + \frac{\lambda}{2} \|x\|^2 \right\},$$

where $\{a_{ij}, b_{ij}\}_{j \in [m]}$ are data samples at the i -th node.

5.1 Data sets and parameter settings

In our experiments we use five standard datasets from the LIBSVM library: **a2a**, **a7a**, **a9a**, **w8a** and **phishing**. Besides, we generated an artificial dataset **artificial** as follows: each of the d elements of the data vector $a_{ij} \in \mathbb{R}^d$ was sampled from the normal distribution $\mathcal{N}(10, 10)$. The corresponding label b_{ij} was sampled uniformly at random from $\{-1, 1\}$. We partitioned each dataset across several nodes (selection of n) in order to capture a variety of scenarios. See Table 3 for more details on all the datasets and the choice of n .

In all experiments we use the theoretical parameters (e.g., stepsizes) for all the three algorithms: vanilla Distributed Compressed Gradient Descent (DCGD) [20], DIANA [29], and ADIANA [23].

Table 3: Data sets used in the experiments, and the number of worker nodes n used in each case.

Data set	# workers n	# data points ($= nm$)	# features d
a2a	15	2 265	123
a7a	100	16 100	123
a9a	80	32 560	123
w8a	142	49 700	300
phishing	100	11 000	68
artificial	100	1 000	200

As the initial approximation of the Hessian in BFGS [6, 9, 11, 43], we use $\mathbf{H}^0 = \nabla^2 P(x^0)$, and the stepsize is 1. We set the same constants in DINGO [7] as they did: $\theta = 10^{-4}$, $\phi = 10^{-6}$, $\rho = 10^{-4}$, and use backtracking line search for DINGO to select the largest stepsize in $\{1, 2^{-1}, 2^{-2}, 2^{-4}, \dots, 2^{-10}\}$. We conduct experiments for three values of the regularization parameter λ : 10^{-3} , 10^{-4} , 10^{-5} . In the figures we plot the relation of the optimality gap $P(x^k) - P(x^*)$ and the number of accumulated transmitted bits⁶ or the number of iterations. The optimal value $P(x^*)$ in each case is the function value at the 20-th iterate of standard Newton’s method. We adopt the realistic setting where the server does not have access to the local data (**Option 1**).

5.2 Compression operators

For the first order methods we use three compression operators: random sparsification [46], random dithering [1], and natural compression [15] (all defined below). For random- r sparsification, the number of communicated bits per iteration is $32r + \log_2 \binom{d}{r}$, and we choose $r = d/4$. For random dithering, we choose $s = \sqrt{d}$, which means the number of communicated bits per iteration is $2.8d + 32$. For natural compression, the number of communicated bits per iteration is $9d$ bits.

For NL1 and NL2 we use the random- r sparsification operator with a selection of values of r . For CNL we use the random sparsification operator \mathcal{C}_p (with $p = 1/20$) induced by the random- r compressor with $r = 1$. This compressor is also defined below.

⁶In all plots, “communicated bits” refers to the total number of bits sent by all nodes to the server.

5.2.1 Random sparsification

The random sparsification compressor [46], denoted random- r , is a randomized mapping $\mathcal{C} : \mathbb{R}^m \rightarrow \mathbb{R}$ defined as

$$\mathcal{C}(x) := \frac{m}{r} \cdot \xi \circ x$$

where $\xi \in \mathbb{R}^m$ is a random vector distributed uniformly at random on the discrete set $\{y \in \{0, 1\}^m : \|y\|_0 = r\}$, where $\|y\|_0 := \{i \mid y_i \neq 0\}$ and \circ is the Hadamard product. The variance parameter associate with this compressor is $\omega = \frac{m}{r} - 1$.

5.2.2 Random dithering

The random dithering compressor [1, 16] with s levels is defined via

$$\mathcal{C}(x) := \text{sign}(x) \cdot \|x\|_q \cdot \frac{1}{s} \cdot \xi_s,$$

where $\|x\|_q := (\sum_i |x_i|^q)^{1/q}$ and $\xi_s \in \mathbb{R}^m$ is a random vector with i -th element defined as

$$\xi_s(i) := \begin{cases} l + 1 & \text{with probability } \frac{|x_i|}{\|x\|_q} s - l \\ l & \text{otherwise} \end{cases}.$$

Here, l satisfies $\frac{|x_i|}{\|x\|_q} \in [\frac{l}{s}, \frac{l+1}{s}]$ and $s \in \mathbb{N}_+$ denotes the levels of the rounding. The variance parameter of this compressor is $\omega \leq 2 + \frac{m^{1/2} + m^{1/q}}{s}$ [16]. For $q = 2$, one can get the improved bound $\omega \leq \min\{\frac{m}{s^2}, \frac{\sqrt{m}}{s}\}$ [1].

5.2.3 Natural compression

The natural compression [15] operator $\mathcal{C}_{\text{nat}} : \mathbb{R}^m \rightarrow \mathbb{R}$ is obtained by applying the random mapping $\mathcal{C} : \mathbb{R} \rightarrow \mathbb{R}$, defined next, to each coordinate of x independently. We define $\mathcal{C}(0) = 0$ and for $t \neq 0$, we let

$$\mathcal{C}(t) := \begin{cases} \text{sign}(t) \cdot 2^{\lceil \log_2 |t| \rceil} & \text{with probability } p(t) := \frac{2^{\lceil \log_2 |t| \rceil} - |t|}{2^{\lceil \log_2 |t| \rceil}} \\ \text{sign}(t) \cdot 2^{\lceil \log_2 |t| \rceil} & \text{with probability } 1 - p(t) \end{cases}$$

The variance parameter of natural compression is $\omega = \frac{1}{8}$.

5.2.4 Bernoulli compressor

A variant of any compression operator $\mathcal{C} : \mathbb{R}^m \rightarrow \mathbb{R}$ can be constructed as follows:

$$\mathcal{C}_p(x) := \begin{cases} \frac{1}{p} \mathcal{C}(x) & \text{with probability } p \\ 0 & \text{otherwise} \end{cases}, \quad (27)$$

where $p \in (0, 1]$ is a probability parameter. It is easy to verify that \mathcal{C}_p is still a compression operator with variance parameter $\omega_p := \frac{\omega+1}{p} - 1$, where ω is the variance parameter of the underlying compressor \mathcal{C} .

5.3 Behavior of NL1 and NL2

Before we compare our methods NL1 and NL2 with competing baselines, we investigate how is their performance affected by the choice of the sparsification parameter r defining the random- r sparsification operator \mathcal{C} . Likewise, we vary the probability parameter p defining the induced Bernoulli compressor \mathcal{C}_p .

According to the results summarized in Figure 1, the best performance of NL1 is obtained for $r = 1$ and $p = 1$. However, for NL2 these parameters are $r = 1$ and $p = 1/20$. We will use these parameter settings for NL1 and NL2 in our subsequent experiments where we compare our methods with several baselines and state-of-the-art methods.

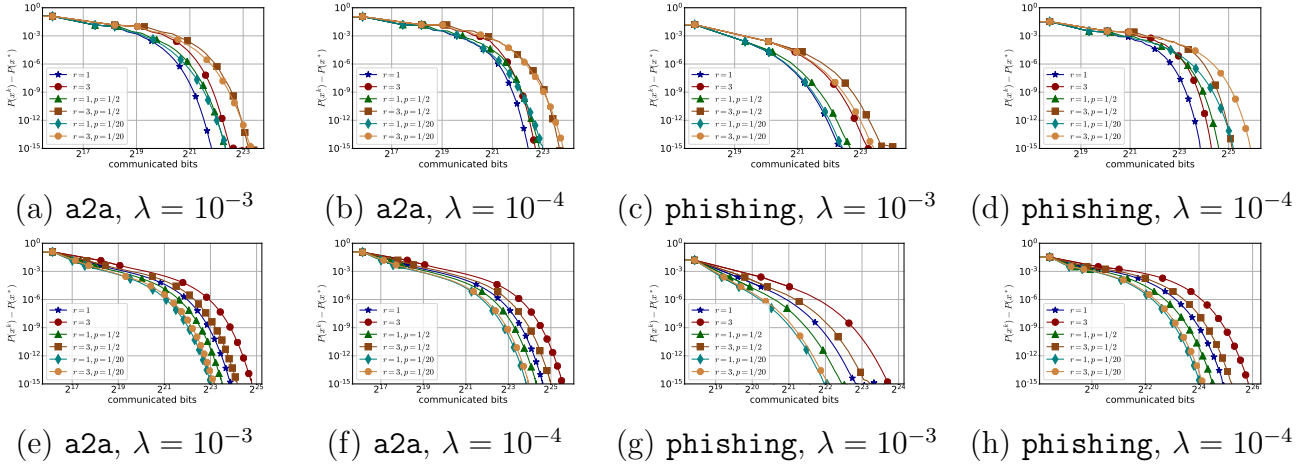


Figure 1: Performance of NL1 (first row) and NL2 (second row) across a few values of r defining the random- r compressor, and a few values of p defining the induced Bernoulli compressor \mathcal{C}_p .

5.4 Comparison of NL1 and NL2 with Newton's method

In our next experiment we compare NL1 and NL2 using different values of r for random- r compression, with Newton's method; see Figure 2. We clearly see that Newton's method performs better than NL1 and NL2 in terms of iteration complexity, as expected. However, our methods have better communication efficiency than Newton's method, by *several orders of magnitude*. Moreover, we see that the smaller r is, the better NL1 and NL2 perform in terms of communication complexity. In Figure 3 we perform a similar comparison for several more datasets, but focus on communication complexity only. The conclusions are unchanged: our methods NL1 and NL2 have superior performance.

5.5 Comparison of NL1 and NL2 with BFGS

In our next test, we compare NL1 and NL2 with BFGS in Figure 4. As we can see, our methods have better communication efficiency than BFGS, by *several orders of magnitude*.

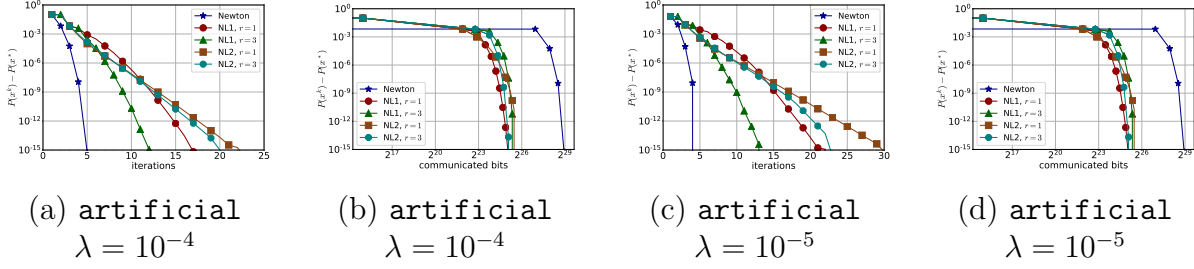


Figure 2: Comparison of NL1 and NL2 with Newton's method in terms of iteration complexity for (a), (c); in terms of communication complexity for (b), (d).

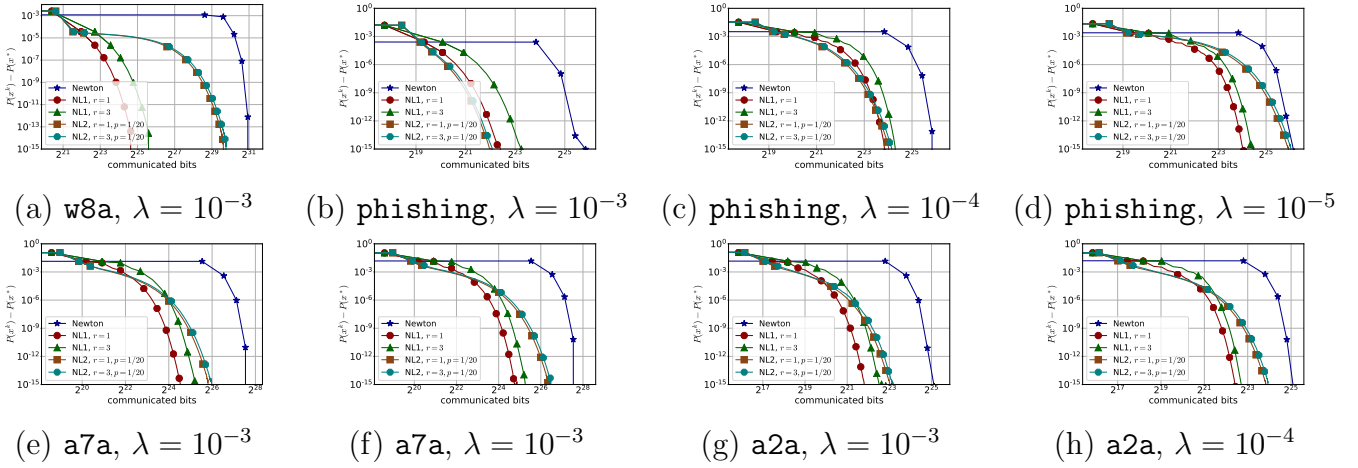


Figure 3: Comparison of NL1, NL2 with Newton's method in terms of communication complexity.

5.6 Comparison of NL1 and NL2 with ADIANA

Next, we compare NL1 and NL2 with ADIANA [23] using three different compression operators: natural compression (DIANA-NC), random sparsification (DIANA-RS, $r = d/4$) and random dithering (DIANA-RD, $s = \sqrt{d}$); see Figure 5. Based on the experimental results, we can conclude that NL1 outperforms all three versions ADIANA for all types of compression, often by *several degrees of magnitude*. The same is true for NL2, with the exception of Figure 5 (c) (w8a dataset) where two variants of ADIANA are faster.

5.7 Comparison of NL1 and NL2 with DINGO

In our next experiment, we compare NL1 and NL2 with DINGO [7]. The results, presented in Figure 6, show that our methods are more communication efficient than DINGO by *many orders of magnitude*. This is true for all experiments with the exception of Figure 6 (c), where DINGO is slightly better than NL2.

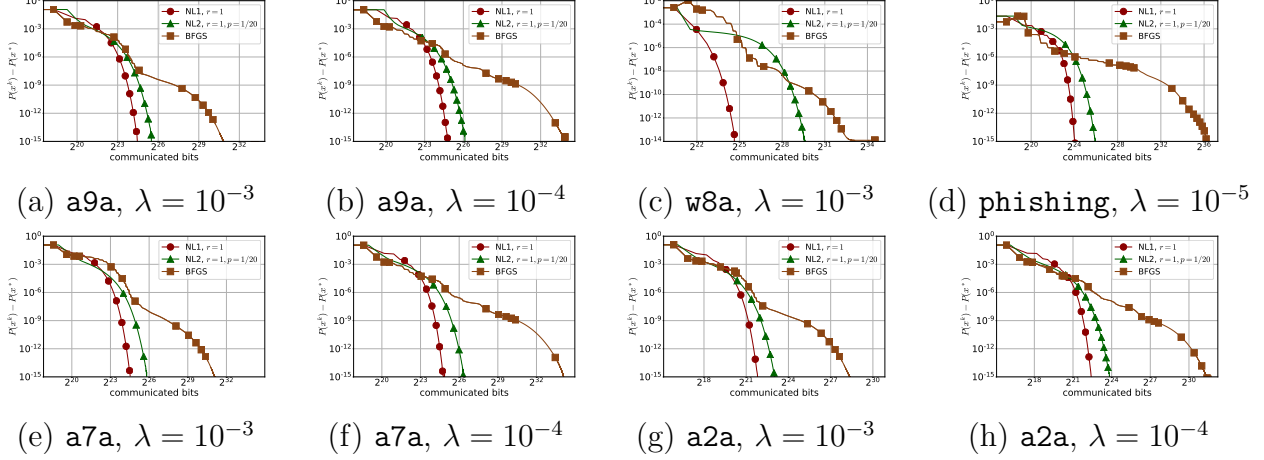


Figure 4: Comparison of NL1, NL2 and BFGS in terms of communication complexity.

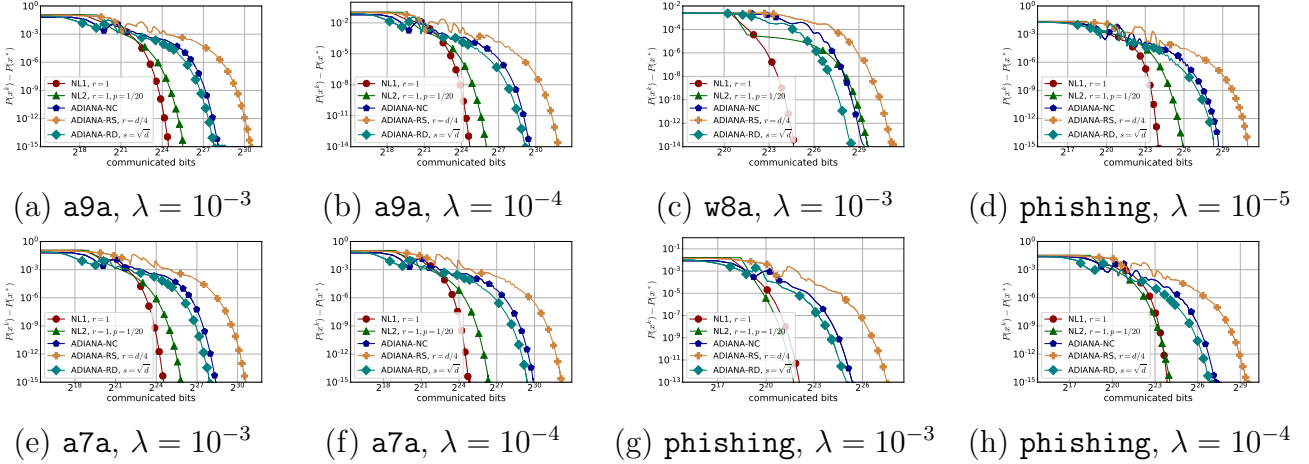


Figure 5: Comparison of NL1, NL2 with ADIANA in terms of communication complexity.

5.8 Comparison of CNL with DCGD and DIANA

We now deploy our CUBIC-NEWTON-LEARN (CNL) method equipped with the random sparsification operator \mathcal{C}_p described in Section 5.2.4. We compare CNL against DIANA [29] and DCGD [20], both in two variants: one using natural compression (NC) and one using random sparsification (RS, $r = d/4$). In Figure 7 we see that for $\lambda = 10^{-3}$, CNL performs slightly better than the gradient-type methods DIANA and DCGD. However, once the regularization parameter becomes small enough ($\lambda = 10^{-4}, 10^{-5}$), the increased condition number hurts the first-order methods, and CNL outperforms both DCGD and DIANA significantly.

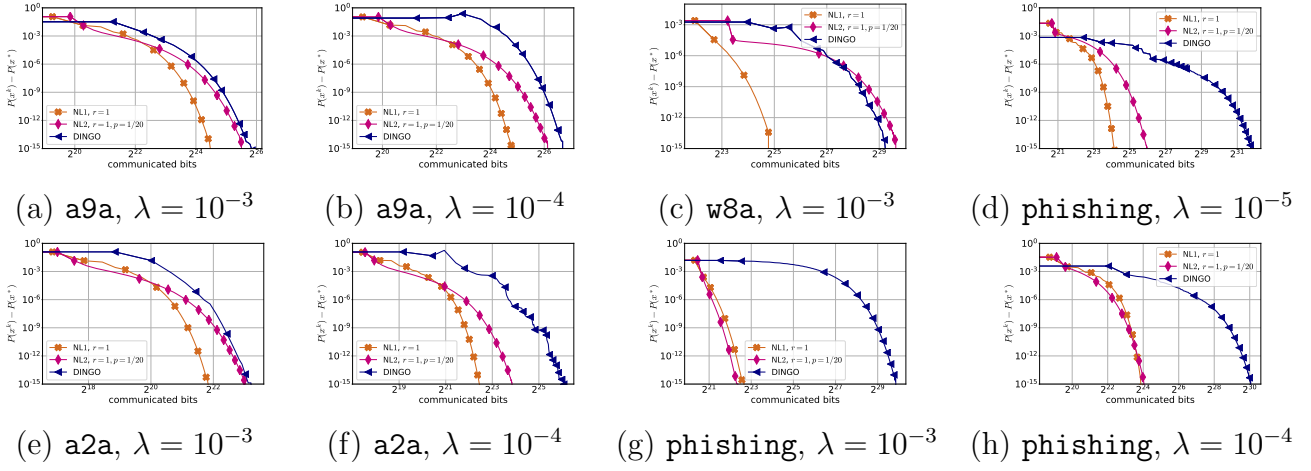


Figure 6: Comparison of NL1, NL2 with ADIANA in terms of communication complexity.

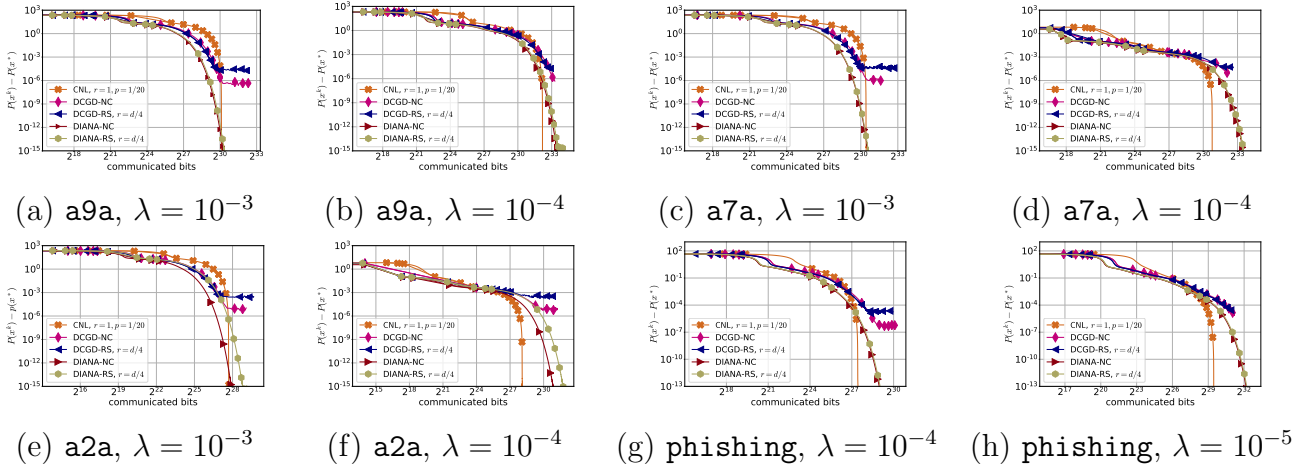


Figure 7: Comparison of CNL with DCGD and DIANA in terms of communication complexity.

Appendix

A Proofs for NL1 (Section 3.5)

Let $W^k := (h_1^k, \dots, h_n^k, x^k)$.

A.1 Lemma

We first establish a lemma.

Lemma A.1. *Let $S^2 := \sum_{i=1}^n \sum_{j=1}^m \|a_{ij}\|^2$. For $\eta \leq \frac{1}{\omega+1}$ and all $k \geq 0$, we have*

$$\mathbb{E} [\mathcal{H}^{k+1} \mid W^k] \leq (1 - \eta) \mathcal{H}^k + \eta \nu^2 S^2 \|x^k - x^*\|^2.$$

Proof. First, recall that

$$\mathcal{H}^k = \sum_{i=1}^n \|h_i^k - h_i(x^*)\|^2. \quad (28)$$

Since $h_i(x^*) \in \mathbb{R}_+^m$ due to convexity of f_{ij} , we have

$$\|[x]_+ - h_i(x^*)\|^2 \leq \|x - h_i(x^*)\|^2, \quad \text{for all } x \in \mathbb{R}^m. \quad (29)$$

Then as long as $\eta \leq \frac{1}{\omega+1}$, we have

$$\begin{aligned} \mathbb{E} [\mathcal{H}^{k+1} \mid W^k] &= \mathbb{E} \left[\sum_{i=1}^n \|h_i^{k+1} - h_i(x^*)\|^2 \mid W^k \right] \\ &= \sum_{i=1}^n \mathbb{E} \left[\|[h_i^k + \eta \mathcal{C}_i^k(h_i(x^k) - h_i^k)]_+ - h_i(x^*)\|^2 \mid W^k \right] \\ &\stackrel{(29)}{\leq} \sum_{i=1}^n \mathbb{E} \left[\|h_i^k + \eta \mathcal{C}_i^k(h_i(x^k) - h_i^k) - h_i(x^*)\|^2 \mid W^k \right] \\ &= \sum_{i=1}^n \|h_i^k - h_i(x^*)\|^2 + 2\eta \sum_{i=1}^n \mathbb{E} [\langle \mathcal{C}_i^k(h_i(x^k) - h_i^k), h_i^k - h_i(x^*) \rangle \mid W^k] \\ &\quad + \eta^2 \sum_{i=1}^n \mathbb{E} [\|\mathcal{C}_i^k(h_i(x^k) - h_i^k)\|^2 \mid W^k] \\ &\stackrel{(28)+(23)+(24)}{\leq} \mathcal{H}^k + 2\eta \sum_{i=1}^n \langle h_i(x^k) - h_i^k, h_i^k - h_i(x^*) \rangle \\ &\quad + \eta^2 \sum_{i=1}^n (\omega + 1) \|h_i(x^k) - h_i^k\|^2. \end{aligned} \quad (30)$$

Using the stepsize restriction $\eta \leq \frac{1}{\omega+1}$, we can bound $\eta^2(\omega + 1) \leq \eta$. Plugging this back in

(30), we get

$$\begin{aligned}
\mathbb{E} [\mathcal{H}^{k+1} \mid W^k] &\leq \mathcal{H}^k + \eta \sum_{i=1}^n \langle h_i(x^k) - h_i^k, h_i(x^k) + h_i^k - 2h_i(x^*) \rangle \\
&= \mathcal{H}^k + \eta \sum_{i=1}^n \left(\|h_i(x^k) - h_i(x^*)\|^2 - \|h_i^k - h_i(x^*)\|^2 \right) \\
&= (1 - \eta) \mathcal{H}^k + \eta \sum_{i=1}^n \|h_i(x^k) - h_i(x^*)\|^2 \\
&\stackrel{(11)}{=} (1 - \eta) \mathcal{H}^k + \eta \sum_{i=1}^n \sum_{j=1}^m (h_{ij}(x^k) - h_{ij}(x^*))^2 \\
&\stackrel{(6)}{\leq} (1 - \eta) \mathcal{H}^k + \eta \sum_{i=1}^n \sum_{j=1}^m \nu^2 \|a_{ij}\|^2 \|x^k - x^*\|^2 \\
&\leq (1 - \eta) \mathcal{H}^k + \eta \nu^2 S^2 \|x^k - x^*\|^2.
\end{aligned}$$

□

A.2 Proof of Theorem 3.2

It is easy to see that

$$\mathbf{H}^k = \frac{1}{n} \sum_{i=1}^n \mathbf{H}_i^k, \quad \mathbf{H}_i^k = \frac{1}{m} \sum_{j=1}^m h_{ij}^k a_{ij} a_{ij}^\top. \quad (31)$$

By the first order necessary optimality conditions, we have

$$\nabla f(x^*) + \lambda x^* = \nabla P(x^*) = 0. \quad (32)$$

Furthermore, since we maintain $\mathbf{H}^k \succeq 0$ for all k , we have $\mathbf{H}^k + \lambda \mathbf{I} \succeq \lambda \mathbf{I}$, whence

$$\|(\mathbf{H}^k + \lambda \mathbf{I})^{-1}\| \leq \frac{1}{\lambda}. \quad (33)$$

Then we can write

$$\begin{aligned}
\|x^{k+1} - x^*\| &\stackrel{(21)}{=} \|x^k - x^* - (\mathbf{H}^k + \lambda \mathbf{I})^{-1}(\nabla f(x^k) + \lambda x^k)\| \\
&= \|(\mathbf{H}^k + \lambda \mathbf{I})^{-1} ((\mathbf{H}^k + \lambda \mathbf{I})(x^k - x^*) - \nabla f(x^k) - \lambda x^k)\| \\
&\stackrel{(33)}{\leq} \frac{1}{\lambda} \|(\mathbf{H}^k + \lambda \mathbf{I})(x^k - x^*) - \nabla f(x^k) - \lambda x^k\| \\
&\stackrel{(32)}{=} \frac{1}{\lambda} \|\mathbf{H}^k(x^k - x^*) - (\nabla f(x^k) - \nabla f(x^*))\| \\
&= \frac{1}{\lambda} \left\| \frac{1}{n} \sum_{i=1}^n [\mathbf{H}_i^k(x^k - x^*) - (\nabla f_i(x^k) - \nabla f_i(x^*))] \right\| \\
&\leq \frac{1}{n\lambda} \sum_{i=1}^n \|\mathbf{H}_i^k(x^k - x^*) - (\nabla f_i(x^k) - \nabla f_i(x^*))\|.
\end{aligned}$$

where the last step follows from applying Jensen's inequality to the convex function $x \mapsto \|x\|$. Using (31) and (2) we get

$$\|x^{k+1} - x^*\| \leq \frac{1}{n\lambda} \sum_{i=1}^n \left\| \frac{1}{m} \sum_{j=1}^m [h_{ij}^k a_{ij} a_{ij}^\top (x^k - x^*) - (\nabla f_{ij}(x^k) - \nabla f_{ij}(x^*))] \right\|. \quad (34)$$

We can now express the difference of the gradients in integral form via the fundamental theorem of calculus, obtaining

$$\begin{aligned} \nabla f_{ij}(x^k) - \nabla f_{ij}(x^*) &= \int_0^1 \mathbf{H}_{ij}(x^* + \tau(x^k - x^*)) (x^k - x^*) d\tau \\ &\stackrel{(4)}{=} \int_0^1 h_{ij}(x^* + \tau(x^k - x^*)) a_{ij} a_{ij}^\top (x^k - x^*) d\tau. \end{aligned}$$

We can plug this back into (34), which gives

$$\begin{aligned} &\|x^{k+1} - x^*\| \\ &\leq \frac{1}{n\lambda} \sum_{i=1}^n \frac{1}{m} \left\| \sum_{j=1}^m \left(h_{ij}^k a_{ij} a_{ij}^\top (x^k - x^*) - \int_0^1 h_{ij}(x^* + \tau(x^k - x^*)) a_{ij} a_{ij}^\top (x^k - x^*) d\tau \right) \right\| \\ &= \frac{1}{n\lambda} \sum_{i=1}^n \frac{1}{m} \left\| \sum_{j=1}^m a_{ij} a_{ij}^\top (x^k - x^*) \left(h_{ij}^k - \int_0^1 h_{ij}(x^* + \tau(x^k - x^*)) d\tau \right) \right\| \\ &\leq \frac{\|x^k - x^*\|}{\lambda} \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m \|a_{ij}\|^2 \left| \int_0^1 h_{ij}^k - h_{ij}(x^* + \tau(x^k - x^*)) d\tau \right|. \quad (35) \end{aligned}$$

From (6), we have

$$\begin{aligned} &\left| \int_0^1 h_{ij}^k - h_{ij}(x^* + \tau(x^k - x^*)) d\tau \right| \\ &\leq \int_0^1 |h_{ij}^k - h_{ij}(x^* + \tau(x^k - x^*))| d\tau \\ &\leq |h_{ij}^k - h_{ij}(x^*)| + \int_0^1 |h_{ij}(x^*) - h_{ij}(x^* + \tau(x^k - x^*))| d\tau \\ &\stackrel{(6)}{\leq} |h_{ij}^k - h_{ij}(x^*)| + \int_0^1 \tau \nu \|a_{ij}\| \cdot \|x^k - x^*\| d\tau \\ &= |h_{ij}^k - h_{ij}(x^*)| + \frac{\nu \|a_{ij}\|}{2} \|x^k - x^*\|. \quad (36) \end{aligned}$$

By squaring both sides of (35), applying Jensen's inequality to the function $t \mapsto t^2$ in the

form $\left(\frac{1}{nm} \sum_i \sum_j t_{ij}\right)^2 \leq \frac{1}{nm} \sum_i \sum_j t_{ij}^2$, and plugging in the bound (36), we get

$$\begin{aligned}
& \|x^{k+1} - x^*\|^2 \\
& \leq \frac{\|x^k - x^*\|^2}{\lambda^2} \left(\frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m \|a_{ij}\|^2 \left| \int_0^1 h_{ij}^k - h_{ij}(x^* + \tau(x^k - x^*)) d\tau \right| \right)^2 \\
& \leq \frac{\|x^k - x^*\|^2}{\lambda^2} \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m \left(\|a_{ij}\|^2 \left| \int_0^1 h_{ij}^k - h_{ij}(x^* + \tau(x^k - x^*)) d\tau \right| \right)^2 \\
& \stackrel{(36)}{\leq} \frac{\|x^k - x^*\|^2}{\lambda^2} \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m \|a_{ij}\|^4 \left(|h_{ij}^k - h_{ij}(x^*)| + \frac{\nu \|a_{ij}\|}{2} \|x^k - x^*\| \right)^2 \\
& \leq \frac{\|x^k - x^*\|^2}{\lambda^2} \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m \|a_{ij}\|^4 \left(2|h_{ij}^k - h_{ij}(x^*)|^2 + \frac{\nu^2 \|a_{ij}\|^2}{2} \|x^k - x^*\|^2 \right) \\
& \leq \frac{\|x^k - x^*\|^2}{\lambda^2} \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m R^4 \left(2|h_{ij}^k - h_{ij}(x^*)|^2 + \frac{\nu^2 R^2}{2} \|x^k - x^*\|^2 \right).
\end{aligned}$$

In the last step we have used Young's inequality.⁷

Since $R := \max_{i,j} \{\|a_{ij}\|\}$, we can further bound

$$\|x^{k+1} - x^*\|^2 \leq \frac{2R^4}{nm\lambda^2} \mathcal{H}^k \|x^k - x^*\|^2 + \frac{\nu^2 R^6}{2\lambda^2} \|x^k - x^*\|^4. \quad (37)$$

Assume $\|x^k - x^*\|^2 \leq \frac{\lambda^2}{12\nu^2 R^6}$ for all $k \geq 0$. Then from (37), we have

$$\begin{aligned}
\|x^{k+1} - x^*\|^2 & \leq \frac{\lambda^2}{12\nu^2 R^6} \cdot \frac{2R^4}{nm\lambda^2} \mathcal{H}^k + \frac{\lambda^2}{12\nu^2 R^6} \cdot \frac{\nu^2 R^6 \|x^k - x^*\|^2}{2\lambda^2} \\
& \leq \frac{1}{6nm\nu^2 R^2} \mathcal{H}^k + \frac{1}{24} \|x^k - x^*\|^2,
\end{aligned}$$

and by taking expectation, we have

$$\mathbb{E} [\|x^{k+1} - x^*\|^2 \mid W^k] \leq \frac{1}{6nm\nu^2 R^2} \mathcal{H}^k + \frac{1}{24} \|x^k - x^*\|^2. \quad (38)$$

Next, Lemma A.1 implies that

$$\mathbb{E} [\mathcal{H}^{k+1} \mid W^k] \leq (1 - \eta) \mathcal{H}^k + \eta nm\nu^2 R^2 \|x^k - x^*\|^2. \quad (39)$$

Recall that $\Phi_1^{k+1} := \|x^{k+1} - x^*\|^2 + \frac{1}{3\eta nm\nu^2 R^2} \mathcal{H}^{k+1}$. Combining (38) and (39), we get

$$\begin{aligned}
\mathbb{E} [\Phi_1^{k+1} \mid W^k] &= \mathbb{E} [\|x^{k+1} - x^*\|^2 \mid W^k] + \frac{1}{3\eta nm\nu^2 R^2} \mathbb{E} [\mathcal{H}^{k+1} \mid W^k] \\
&\stackrel{(38)}{\leq} \frac{1}{3\eta nm\nu^2 R^2} \left(1 - \eta + \frac{\eta}{2} \right) \mathcal{H}^k + \left(\frac{1}{24} + \frac{1}{3} \right) \|x^k - x^*\|^2 \\
&\leq \left(1 - \min \left\{ \frac{\eta}{2}, \frac{5}{8} \right\} \right) \Phi_1^k = \theta_1^k \Phi_1^k.
\end{aligned} \quad (40)$$

⁷ $(a+b)^2 \leq 2a^2 + 2b^2$.

By applying the tower property, we get

$$\mathbb{E} [\Phi_1^{k+1}] = \mathbb{E} [\mathbb{E} [\Phi_1^{k+1} | W^k]] \stackrel{(40)}{\leq} \theta_1 \mathbb{E} [\Phi_1^k].$$

Unrolling the recursion, we get $\mathbb{E} [\Phi_1^k] \leq \theta_1^k \Phi_1^0$, and the first claim is proved.

We further have $\mathbb{E} [\|x^k - x^*\|^2] \leq \theta_1^k \Phi_1^0$ and $\mathbb{E} [\mathcal{H}^k] \leq \theta_1^k 3\eta nm \nu^2 R^2 \Phi_1^0$. Assume $x^k \neq x^*$ for all k . Then from (37), we have

$$\frac{\|x^{k+1} - x^*\|^2}{\|x^k - x^*\|^2} \leq \frac{2R^4}{nm\lambda^2} \mathcal{H}^k + \frac{\nu^2 R^6}{2\lambda^2} \|x^k - x^*\|^2,$$

and by taking expectation, we obtain

$$\begin{aligned} \mathbb{E} \left[\frac{\|x^{k+1} - x^*\|^2}{\|x^k - x^*\|^2} \right] &\leq \frac{2R^4}{mn\lambda^2} \mathbb{E} [\mathcal{H}^k] + \frac{\nu^2 R^6}{2\lambda^2} \mathbb{E} [\|x^k - x^*\|^2] \\ &\leq \theta_1^k \left(6\eta + \frac{1}{2} \right) \frac{\nu^2 R^6}{\lambda^2} \Phi_1^0. \end{aligned}$$

A.3 Proof of Lemma 3.3

We prove this by induction. First, we have $\|x^0 - x^*\|^2 \leq \frac{\lambda^2}{12\nu^2 R^6}$ by the assumption. We assume $\|x^k - x^*\|^2 \leq \frac{\lambda^2}{12\nu^2 R^6}$ holds for all $k \leq K$. For $k \leq K$, since h_{ij}^k is a convex combination of $\{h_{ij}(x^0), h_{ij}(x^1), \dots, h_{ij}(x^k)\}$ for all i, j , we have

$$h_{ij}^k = \sum_{t=0}^k \rho_t h_{ij}(x^t), \quad \text{with} \quad \sum_{t=0}^k \rho_t = 1, \quad \text{and} \quad \rho_t \geq 0,$$

which implies that

$$\begin{aligned} \sum_{i=1}^n \|h_i^k - h_i(x^*)\|^2 &= \sum_{i=1}^n \sum_{j=1}^m |h_{ij}^k - h_{ij}(x^*)|^2 \\ &= \sum_{i=1}^n \sum_{j=1}^m \left| \sum_{t=0}^k \rho_t (h_{ij}(x^t) - h_{ij}(x^*)) \right|^2 \\ &\leq \sum_{i=1}^n \sum_{j=1}^m \sum_{t=0}^k \rho_t |h_{ij}(x^t) - h_{ij}(x^*)|^2 \\ &\stackrel{(6)}{\leq} \sum_{i=1}^n \sum_{j=1}^m \sum_{t=0}^k \rho_t \nu^2 \|a_{ij}\|^2 \|x^t - x^*\|^2 \\ &\stackrel{\text{As. 3.1}}{\leq} \nu^2 R^2 \sum_{i=1}^n \sum_{j=1}^m \sum_{t=0}^k \rho_t \cdot \frac{\lambda^2}{12\nu^2 R^6} \\ &= \frac{mn\lambda^2}{12R^4}, \end{aligned}$$

for $k \leq K$.

Combining the above inequality and (37), we arrive at

$$\begin{aligned}
\|x^{K+1} - x^*\|^2 &\leq \frac{2\|x^K - x^*\|^2 R^4}{mn\lambda^2} \sum_{i=1}^n \|h_i^K - h_i(x^*)\|^2 + \frac{\nu^2 R^6 \|x^K - x^*\|^4}{2\lambda^2} \\
&\leq \frac{2\|x^K - x^*\|^2 R^4}{mn\lambda^2} \cdot \frac{mn\lambda^2}{12R^4} + \frac{\nu^2 R^6 \|x^K - x^*\|^4}{2\lambda^2} \\
&\leq \frac{1}{6} \|x^K - x^*\|^2 + \frac{\nu^2 R^6 \|x^K - x^*\|^4}{2\lambda^2} \\
&\leq \frac{1}{6} \cdot \frac{\lambda^2}{12\nu^2 R^6} + \frac{\nu^2 R^6}{2\lambda^2} \cdot \left(\frac{\lambda^2}{12\nu^2 R^6} \right)^2 \\
&\leq \frac{\lambda^2}{12\nu^2 R^6}.
\end{aligned}$$

B Proofs for NL2 (Section 3.6)

For $\mathcal{H}^k := \sum_{i=1}^n \|h_i^k - h_i(x^*)\|^2$, even though the update of h_i^k in Algorithm 2 is slightly different from that of Algorithm 1, we also have the following lemma. The proof is almost the same as that of Lemma A.1, hence we omit it.

Lemma B.1. *For $\eta \leq \frac{1}{\omega+1}$, we have*

$$\mathbb{E}[\mathcal{H}^{k+1}] \leq (1 - \eta)\mathbb{E}[\mathcal{H}^k] + mn\eta\nu^2 R^2 \mathbb{E}[\|x^k - x^*\|^2],$$

for $k \geq 0$.

B.1 Proof of Theorem 3.5

First, we have

$$\begin{aligned} \|x^{k+1} - x^*\| &= \|x^k - x^* - (\mathbf{H}^k + \lambda\mathbf{I})^{-1}(\nabla f(x^k) + \lambda x^k)\| \\ &= \|(\mathbf{H}^k + \lambda\mathbf{I})^{-1}((\mathbf{H}^k + \lambda\mathbf{I})(x^k - x^*) - \nabla f(x^k) - \lambda x^k)\| \\ &\leq \frac{1}{\mu} \|(\mathbf{H}^k + \lambda\mathbf{I})(x^k - x^*) - \nabla f(x^k) - \lambda x^k\| \\ &\stackrel{(32)}{=} \frac{1}{\mu} \|\mathbf{H}^k(x^k - x^*) - (\nabla f(x^k) - \nabla f(x^*))\| \\ &\leq \frac{1}{n\mu} \sum_{i=1}^n \|\mathbf{H}_i^k(x^k - x^*) - (\nabla f_i(x^k) - \nabla f_i(x^*))\| \\ &= \frac{1}{n\mu} \sum_{i=1}^n \frac{1}{m} \left\| \sum_{j=1}^m ((\beta^k(h_{ij}^k + 2\gamma) - 2\gamma) a_{ij} a_{ij}^\top (x^k - x^*) \right. \\ &\quad \left. - (\nabla f_{ij}(x^k) - \nabla f_{ij}(x^*))) \right\| \end{aligned}$$

By using (17), we further get

$$\begin{aligned} &\|x^{k+1} - x^*\| \\ &\stackrel{(17)}{=} \frac{1}{n\mu} \sum_{i=1}^n \frac{1}{m} \left\| \sum_{j=1}^m (\beta^k(h_{ij}^k + 2\gamma) a_{ij} a_{ij}^\top (x^k - x^*) - 2\gamma a_{ij} a_{ij}^\top (x^k - x^*) \right. \\ &\quad \left. - \int_0^1 \mathbf{H}_{ij}(x^* + \tau(x^k - x^*))(x^k - x^*) d\tau) \right\| \\ &= \frac{1}{n\mu} \sum_{i=1}^n \frac{1}{m} \left\| \sum_{j=1}^m (\beta^k(h_{ij}^k + 2\gamma) a_{ij} a_{ij}^\top (x^k - x^*) - 2\gamma a_{ij} a_{ij}^\top (x^k - x^*) \right. \\ &\quad \left. - \int_0^1 h_{ij}(x^* + \tau(x^k - x^*)) a_{ij} a_{ij}^\top (x^k - x^*) d\tau) \right\| \\ &= \frac{1}{n\mu} \sum_{i=1}^n \frac{1}{m} \left\| \sum_{j=1}^m (h_{ij}^k + 2\gamma) a_{ij} a_{ij}^\top (x^k - x^*) \left(\beta^k - \int_0^1 \frac{h_{ij}(x^* + \tau(x^k - x^*)) + 2\gamma}{h_{ij}^k + 2\gamma} d\tau \right) \right\| \\ &\leq \frac{3\gamma \|x^k - x^*\|}{n\mu} \sum_{i=1}^n \sum_{j=1}^m \frac{\|a_{ij}\|^2}{m} \left| \beta^k - \int_0^1 \frac{h_{ij}(x^* + \tau(x^k - x^*)) + 2\gamma}{h_{ij}^k + 2\gamma} d\tau \right|, \end{aligned}$$

where we use $|h_{ij}^k| \leq \gamma$ in the last inequality. Next, we estimate $|\beta^k - 1|$:

$$\begin{aligned}
\left| \frac{h_{ij}(x^k) + 2\gamma}{h_{ij}^k + 2\gamma} - 1 \right| &= \left| \frac{h_{ij}(x^k) - h_{ij}^k}{h_{ij}^k + 2\gamma} \right| \\
&\leq \frac{1}{\gamma} |h_{ij}(x^k) - h_{ij}(x^*) + h_{ij}(x^*) - h_{ij}^k| \\
&\leq \frac{1}{\gamma} |h_{ij}(x^k) - h_{ij}(x^*)| + \frac{1}{\gamma} |h_{ij}^k - h_{ij}(x^*)| \\
&\stackrel{(6)}{\leq} \frac{\nu \|a_{ij}\|}{\gamma} \|x^k - x^*\| + \frac{1}{\gamma} |h_{ij}^k - h_{ij}(x^*)|.
\end{aligned}$$

Let $\{i^k, j^k\} = \operatorname{argmax}_{i,j} \left\{ \frac{h_{ij}(x^k) + 2\gamma}{h_{ij}^k + 2\gamma} \right\}$. Then we have

$$\begin{aligned}
|\beta^k - 1| &= \left| \frac{h_{i^k j^k}(x^k) + 2\gamma}{h_{i^k j^k}^k + 2\gamma} - 1 \right| \\
&\leq \max_{i,j} \left\{ \frac{\nu \|a_{ij}\|}{\gamma} \|x^k - x^*\| + \frac{1}{\gamma} |h_{ij}^k - h_{ij}(x^*)| \right\} \\
&\leq \frac{\nu R}{\gamma} \|x^k - x^*\| + \frac{1}{\gamma} \max_{i,j} \{|h_{ij}^k - h_{ij}(x^*)|\}.
\end{aligned}$$

For $\left| \int_0^1 \frac{h_{ij}(x^* + \tau(x^k - x^*)) + 2\gamma}{h_{ij}^k + 2\gamma} d\tau - 1 \right|$, we have

$$\begin{aligned}
&\left| \int_0^1 \frac{h_{ij}(x^* + \tau(x^k - x^*)) + 2\gamma}{h_{ij}^k + 2\gamma} d\tau - 1 \right| \\
&= \left| \int_0^1 \frac{h_{ij}(x^* + \tau(x^k - x^*)) - h_{ij}^k}{h_{ij}^k + 2\gamma} d\tau \right| \\
&\leq \int_0^1 \left| \frac{h_{ij}(x^* + \tau(x^k - x^*)) - h_{ij}^k}{h_{ij}^k + 2\gamma} \right| d\tau \\
&\leq \frac{1}{\gamma} \int_0^1 (|h_{ij}(x^* + \tau(x^k - x^*)) - h_{ij}(x^*)| + |h_{ij}^k - h_{ij}(x^*)|) d\tau \\
&\stackrel{(6)}{\leq} \frac{1}{\gamma} |h_{ij}^k - h_{ij}(x^*)| + \frac{1}{\gamma} \int_0^1 \tau \nu \|a_{ij}\| \cdot \|x^k - x^*\| d\tau \\
&= \frac{1}{\gamma} |h_{ij}^k - h_{ij}(x^*)| + \frac{\nu \|a_{ij}\|}{2\gamma} \|x^k - x^*\|.
\end{aligned}$$

Combining the two above inequalities, we can obtain

$$\begin{aligned}
&\left| \beta^k - \int_0^1 \frac{h_{ij}(x^* + \tau(x^k - x^*)) + 2\gamma}{h_{ij}^k + 2\gamma} d\tau \right| \\
&\leq \frac{\nu R}{\gamma} \|x^k - x^*\| + \frac{1}{\gamma} \max_{i,j} \{|h_{ij}^k - h_{ij}(x^*)|\} + \frac{1}{\gamma} |h_{ij}^k - h_{ij}(x^*)| + \frac{\nu \|a_{ij}\|}{2\gamma} \|x^k - x^*\| \\
&\leq \frac{2\nu R}{\gamma} \|x^k - x^*\| + \frac{2}{\gamma} \max_{i,j} \{|h_{ij}^k - h_{ij}(x^*)|\}. \tag{41}
\end{aligned}$$

Thus, we have

$$\begin{aligned}
& \|x^{k+1} - x^*\|^2 \\
& \leq \frac{9\gamma^2 \|x^k - x^*\|^2}{n\mu^2} \sum_{i=1}^n \sum_{j=1}^m \frac{\|a_{ij}\|^4}{m} \left| \beta^k - \int_0^1 \frac{h_{ij}(x^* + \tau(x^k - x^*)) + 2\gamma}{h_{ij}^k + 2\gamma} d\tau \right|^2 \\
& \leq \frac{9\gamma^2 \|x^k - x^*\|^2}{n\mu^2} \sum_{i=1}^n \sum_{j=1}^m \frac{\|a_{ij}\|^4}{m} \left(\frac{8\nu^2 R^2}{\gamma^2} \|x^k - x^*\|^2 + \frac{8}{\gamma^2} \max_{i,j} \{|h_{ij}^k - h_{ij}(x^*)|^2\} \right) \\
& \leq \frac{72\|x^k - x^*\|^2 R^4}{\mu^2} \max_{i,j} \{|h_{ij}^k - h_{ij}(x^*)|^2\} + \frac{72\nu^2 R^6 \|x^k - x^*\|^4}{\mu^2} \\
& \leq \frac{72\|x^k - x^*\|^2 R^4}{\mu^2} \sum_{i=1}^n \|h_i^k - h_i(x^*)\|^2 + \frac{72\nu^2 R^6 \|x^k - x^*\|^4}{\mu^2},
\end{aligned} \tag{42}$$

where we use the convexity of $x \mapsto \|x\|^2$ in the first inequality. From the definition of \mathcal{H}^k , we can write

$$\|x^{k+1} - x^*\|^2 \leq \frac{72\|x^k - x^*\|^2 R^4}{\mu^2} \mathcal{H}^k + \frac{72\nu^2 R^6 \|x^k - x^*\|^4}{\mu^2}. \tag{43}$$

Assume $\|x^k - x^*\|^2 \leq \frac{\mu^2}{432mn\nu^2 R^6}$ for all $k \geq 0$. Then from (43), we have

$$\begin{aligned}
\|x^{k+1} - x^*\|^2 & \leq \frac{\mu^2}{432mn\nu^2 R^6} \cdot \frac{72R^4}{\mu^2} \mathcal{H}^k + \frac{\mu^2}{432mn\nu^2 R^6} \cdot \frac{72\nu^2 R^6 \|x^k - x^*\|^2}{\mu^2} \\
& \leq \frac{1}{6mn\nu^2 R^2} \mathcal{H}^k + \frac{1}{6mn} \|x^k - x^*\|^2,
\end{aligned}$$

and by taking expectation, we have

$$\mathbb{E} [\|x^{k+1} - x^*\|^2] \leq \frac{1}{6mn\nu^2 R^2} \mathbb{E} [\mathcal{H}^k] + \frac{1}{6mn} \mathbb{E} [\|x^k - x^*\|^2]. \tag{44}$$

Let $\Phi_2^k := \|x^k - x^*\|^2 + \frac{1}{3mn\nu^2 R^2} \mathcal{H}^k$. Combining (44) and the evolution of \mathcal{H}^k in Lemma B.1, we arrive at

$$\begin{aligned}
\mathbb{E}[\Phi_2^{k+1}] & = \mathbb{E} [\|x^{k+1} - x^*\|^2] + \frac{1}{3mn\nu^2 R^2} \mathbb{E} [\mathcal{H}^{k+1}] \\
& \leq \frac{1}{3mn\nu^2 R^2} \left(1 - \eta + \frac{\eta}{2} \right) \mathbb{E}[\mathcal{H}^k] + \left(\frac{1}{6mn} + \frac{1}{3} \right) \mathbb{E} [\|x^k - x^*\|^2] \\
& \leq \left(1 - \min \left\{ \frac{\eta}{2}, \frac{1}{2} \right\} \right) \mathbb{E} [\Phi_2^k],
\end{aligned}$$

which implies that $\mathbb{E} [\Phi_2^k] \leq \left(1 - \min \left\{ \frac{\eta}{2}, \frac{1}{2} \right\} \right)^k \Phi_2^0$. Then we further have $\mathbb{E} [\|x^k - x^*\|^2] \leq \left(1 - \min \left\{ \frac{\eta}{2}, \frac{1}{2} \right\} \right)^k \Phi_2^0$ and $\mathbb{E}[\mathcal{H}^k] \leq \left(1 - \min \left\{ \frac{\eta}{2}, \frac{1}{2} \right\} \right)^k 3mn\nu^2 R^2 \Phi_2^0$. Assume $x^k \neq x^*$ for all k . Then from (43), we have

$$\frac{\|x^{k+1} - x^*\|^2}{\|x^k - x^*\|^2} \leq \frac{72R^4}{\mu^2} \mathcal{H}^k + \frac{72\nu^2 R^6}{\mu^2} \|x^k - x^*\|^2,$$

and by taking expectation, we can get

$$\begin{aligned}\mathbb{E} \left[\frac{\|x^{k+1} - x^*\|^2}{\|x^k - x^*\|^2} \right] &\leq \frac{72R^4}{\mu^2} \mathbb{E} [\mathcal{H}^k] + \frac{72\nu^2 R^6}{\mu^2} \mathbb{E} [\|x^k - x^*\|^2] \\ &\leq \left(1 - \min \left\{ \frac{\eta}{2}, \frac{1}{2} \right\} \right)^k (3mn\eta + 1) \frac{72\nu^2 R^6}{\mu^2} \Phi_2^0.\end{aligned}$$

B.2 Proof of Lemma 3.6

We prove this by induction. First, we have $\|x^0 - x^*\|^2 \leq \frac{\mu^2}{432mn\nu^2 R^6}$ by the assumption. We assume $\|x^k - x^*\|^2 \leq \frac{\mu^2}{432mn\nu^2 R^6}$ holds for all $k \leq K$. For $k \leq K$, since h_{ij}^k is a convex combination of $\{h_{ij}(x^0), h_{ij}(x^1), \dots, h_{ij}(x^k)\}$ for all i, j , we have

$$h_{ij}^k = \sum_{t=0}^k \rho_t h_{ij}(x^t), \quad \text{with} \quad \sum_{t=0}^k \rho_t = 1, \quad \text{and} \quad \rho_t \geq 0,$$

which implies that

$$\begin{aligned}|h_{ij}^k - h_{ij}(x^*)|^2 &= \left| \sum_{t=0}^k \rho_t (h_{ij}(x^t) - h_{ij}(x^*)) \right|^2 \\ &\leq \sum_{t=0}^k \rho_t |h_{ij}(x^t) - h_{ij}(x^*)|^2 \\ &\stackrel{(6)}{\leq} \sum_{t=0}^k \rho_t \nu^2 \|a_{ij}\|^2 \|x^t - x^*\|^2 \\ &\stackrel{\text{As.3.4}}{\leq} \nu^2 R^2 \sum_{t=0}^k \rho_t \cdot \frac{\mu^2}{432mn\nu^2 R^6} \\ &= \frac{\mu^2}{432mnR^4},\end{aligned}$$

for $k \leq K$.

Combining the above inequality and (42), we arrive at

$$\begin{aligned}\|x^{K+1} - x^*\|^2 &\leq \frac{72\|x^K - x^*\|^2 R^4}{\mu^2} \max_{i,j} \{|h_{ij}^K - h_{ij}(x^*)|^2\} + \frac{72\nu^2 R^6 \|x^K - x^*\|^4}{\mu^2} \\ &\leq \frac{72\|x^K - x^*\|^2 R^4}{\mu^2} \cdot \frac{\mu^2}{432mnR^4} + \frac{72\nu^2 R^6 \|x^K - x^*\|^4}{\mu^2} \\ &\leq \frac{1}{6mn} \|x^K - x^*\|^2 + \frac{72\nu^2 R^6 \|x^K - x^*\|^4}{\mu^2} \\ &\leq \frac{1}{6mn} \cdot \frac{\mu^2}{432\nu^2 R^6} + \frac{72\nu^2 R^6}{\mu^2} \cdot \left(\frac{\mu^2}{432mn\nu^2 R^6} \right)^2 \\ &\leq \frac{\mu^2}{432mn\nu^2 R^6}.\end{aligned}$$

C Proofs for CNL (Section 4)

C.1 Solving the Subproblem

In CUBIC-NEWTON-LEARN, we need to minimize $T(x^k, s)$ at each step. The optimality condition for this subproblem is

$$\nabla P(x^k) + (\mathbf{H}^k + \lambda \mathbf{I})s + \frac{M\|s\|}{2}s = 0. \quad (45)$$

Let $\mathbf{U}^T \Lambda \mathbf{U}$ be the eigenvalue decomposition of $\mathbf{H}^k + \lambda \mathbf{I}$. Then we can transform the above equality to

$$\Lambda \mathbf{U}s + \frac{M\|s\|}{2}\mathbf{U}s = -\mathbf{U}\nabla P(x^k).$$

Noticing that $\|\mathbf{U}s\| = \|s\|$, define $y = \mathbf{U}s$, then we can get

$$y = -\left(\Lambda + \frac{M}{2}\|y\|\mathbf{I}\right)^{-1} \mathbf{U}\nabla P(x^k),$$

and by taking the norm, we arrive at

$$\|y\|^2 = \sum_{i=1}^d \frac{(\mathbf{U}\nabla P(x^k))_i^2}{(\Lambda_i + \frac{M}{2}\|y\|)^2},$$

which is actually equivalent to a one-dimensional nonlinear equation. By solving this one-dimensional nonlinear equation, we can get $\|s\| = \|\mathbf{U}s\| = \|y\|$, and then from (45), we can obtain the solution

$$s = -\left(\mathbf{H}^k + \lambda \mathbf{I} + \frac{M\|s\|\mathbf{I}}{2}\right)^{-1} \nabla P(x^k) = -\left(\mathbf{H}^k + \lambda \mathbf{I} + \frac{M\|y\|\mathbf{I}}{2}\right)^{-1} \nabla P(x^k).$$

C.2 Proof of Lemma 4.2

Since $x^{k+1} = x^k + s^k$, from Lemma 1 in [33], we have

$$\begin{aligned} P(x^{k+1}) &\leq P(x^k) + \langle \nabla f(x^k) + \lambda x^k, s^k \rangle + \frac{1}{2} \langle (\nabla^2 f(x^k) + \lambda \mathbf{I})s^k, s^k \rangle + \frac{M}{6} \|s^k\|^3 \\ &\leq P(x^k) + \langle \nabla f(x^k) + \lambda x^k, s^k \rangle + \frac{1}{2} \langle (\mathbf{H}^k + \lambda \mathbf{I})s^k, s^k \rangle + \frac{M}{6} \|s^k\|^3 \\ &= P(x^k) + T(x^k, s^k). \end{aligned} \quad (46)$$

Since $T(x^k, s^k) = \min_s \{T(x^k, s)\} \leq T(x^k, 0) = 0$, we have $P(x^{k+1}) \leq P(x^k)$. From (46), we have

$$\begin{aligned} P(x^{k+1}) &\leq P(x^k) + T(x^k, s^k) \\ &= P(x^k) + \min_s \{T(x^k, s)\} \\ &\leq P(x^k) + T(x^k, y - x^k) \\ &= P(x^k) + \langle \nabla P(x^k), y - x^k \rangle + \frac{1}{2} \langle (\mathbf{H}^k + \lambda \mathbf{I})(y - x^k), y - x^k \rangle + \frac{M}{6} \|y - x^k\|^2 \\ &\leq P(x^k) + \langle \nabla P(x^k), y - x^k \rangle + \frac{1}{2} \langle \nabla^2 P(x^k)(y - x^k), y - x^k \rangle + \frac{M}{6} \|y - x^k\|^2 \\ &\quad + \frac{1}{2} \|\mathbf{H}^k + \lambda \mathbf{I} - \nabla^2 P(x^k)\| \cdot \|y - x^k\|^2. \end{aligned}$$

First, by Lemma 1 in [33], we can obtain

$$\langle \nabla P(x^k), y - x^k \rangle + \frac{1}{2} \langle \nabla^2 P(x^k)(y - x^k), y - x^k \rangle \leq P(y) - P(x^k) + \frac{M}{6} \|y - x^k\|^3.$$

For $\|\mathbf{H}^k + \lambda \mathbf{I} - \nabla^2 P(x^k)\|$, we have

$$\begin{aligned} \|\mathbf{H}^k + \lambda \mathbf{I} - \nabla^2 P(x^k)\| &= \|\mathbf{H}^k - \nabla^2 f(x^k)\| \\ &= \left\| \frac{1}{n} \sum_{i=1}^n \frac{1}{m} \sum_{j=1}^m (\beta^k(h_{ij}^k + 2\gamma) - 2\gamma - h_{ij}(x^k)) a_{ij} a_{ij}^\top \right\| \\ &\leq \frac{1}{n} \sum_{i=1}^n \frac{1}{m} \sum_{j=1}^m R^2 |\beta^k(h_{ij}^k + 2\gamma) - 2\gamma - h_{ij}(x^k)| \\ &\leq \frac{1}{n} \sum_{i=1}^n \frac{1}{m} \sum_{j=1}^m R^2 |h_{ij}^k + 2\gamma| \cdot \left| \beta^k - \frac{h_{ij}(x^k) + 2\gamma}{h_{ij}^k + 2\gamma} \right|. \end{aligned}$$

Since $|h_{ij}^k| \leq \gamma$ and $|h_{ij}(x^k)| \leq \gamma$, we have $|h_{ij}^k + 2\gamma| \leq 3\gamma$ and $\left| \frac{h_{ij}(x^k) + 2\gamma}{h_{ij}^k + 2\gamma} \right| \leq 3$ for any i, j . Thus we can get

$$\|\mathbf{H}^k + \lambda \mathbf{I} - \nabla^2 P(x^k)\| \leq 18\gamma R^2.$$

Then we arrive at

$$P(x^{k+1}) \leq P(y) + 9\gamma R^2 \|y - x^k\|^2 + \frac{M}{3} \|y - x^k\|^3.$$

C.3 Proof of Theorem 4.3

Recall that $R_x = \sup_{x \in \mathbb{R}^d} \{\|x - x^*\| : P(x) \leq P(x^0)\}$. Since $P(x^k) \leq P(x^0)$ for all $k \geq 0$, we have $\|x^k - x^*\| \leq R_x$ for $k \geq 0$. Let $c_k := k^2$, $C_k = C_0 + \sum_{i=1}^k c_i$ with $C_0 = \frac{4}{3}$ for $k \geq 1$. Then we can obtain

$$C_k = C_0 + \sum_{i=1}^k i^2 \geq C_0 + \int_0^k x^2 dx = C_0 + \frac{k^3}{3}. \quad (47)$$

Let $\sigma_k := \frac{c_{k+1}}{C_{k+1}}$ for $k \geq 0$. Then $1 - \sigma_k = \frac{C_k}{C_{k+1}}$. We also have $\sigma_k \leq 1$ since

$$\max_{\xi \geq 0} \frac{\xi^2}{C_0 + \frac{\xi^3}{3}} = 1.$$

Choose $y = \sigma_k x^* + (1 - \sigma_k) x^k$. Then by the convexity of P , we have

$$\begin{aligned} P(x^{k+1}) &\leq \sigma_k P(x^*) + (1 - \sigma_k) P(x^k) + 9\gamma R^2 \sigma_k^2 \|x^k - x^*\|^2 + \frac{M\sigma_k^3}{3} \|x^k - x^*\|^3 \\ &\leq \frac{c_{k+1}}{C_{k+1}} P(x^*) + \frac{C_k}{C_{k+1}} P(x^k) + 9\gamma R^2 R_x^2 \left(\frac{c_{k+1}}{C_{k+1}} \right)^2 + \frac{MR_x^3}{3} \left(\frac{c_{k+1}}{C_{k+1}} \right)^3, \end{aligned}$$

which implies that

$$C_{k+1}[P(x^{k+1}) - P(x^*)] \leq C_k[P(x^k) - P(x^*)] + 9\gamma R^2 R_x^2 \frac{c_{k+1}^2}{C_{k+1}} + \frac{MR_x^3}{3} \frac{c_{k+1}^3}{C_{k+1}^2}.$$

Hence,

$$C_k[P(x^k) - P(x^*)] \leq C_0[P(x^0) - P(x^*)] + 9\gamma R^2 R_x^2 \sum_{i=1}^k \frac{c_i^2}{C_i} + \frac{MR_x^3}{3} \sum_{i=1}^k \frac{c_i^3}{C_i^2}.$$

Notice that

$$\sum_{i=1}^k \frac{c_i^2}{C_i} \stackrel{(47)}{\leq} \sum_{i=1}^k \frac{i^4}{C_0 + \frac{1}{3}i^3} \leq 3 \sum_{i=1}^k i \leq 3k^2,$$

and

$$\sum_{i=1}^k \frac{c_i^3}{C_i^2} \stackrel{(47)}{\leq} \sum_{i=1}^k \frac{i^6}{(C_0 + \frac{1}{3}i^3)^2} \leq 9k.$$

Therefore, we arrive at

$$\begin{aligned} P(x^k) - P(x^*) &\leq \frac{1}{C_k} (C_0[P(x^0) - P(x^*)] + 27\gamma R^2 R_x^2 k^2 + 3MR_x^3 k) \\ &\stackrel{(47)}{\leq} \frac{81\gamma R^2 R_x^2}{k} + \frac{9MR_x^3}{k^2} + \frac{4[P(x^0) - P(x^*)]}{k^3}. \end{aligned}$$

Then in order to guarantee $P(x^k) - P(x^*) \leq \epsilon$, we only need to let

$$k = O\left(\frac{81\gamma R^2 R_x^2}{\epsilon} + \sqrt{\frac{MR_x^3}{\epsilon}} + \left(\frac{P(x^0) - P(x^*)}{\epsilon}\right)^{1/3}\right).$$

C.4 Proof of Theorem 4.4

Since P is μ -strongly convex, we have $\frac{\mu}{2}\|x - x^*\|^2 \leq P(x) - P(x^*)$, and thus

$$R_x^2 \leq \frac{2}{\mu}(P(x^0) - P(x^*)).$$

Combining the above inequality and Theorem 4.3, we can obtain

$$P(x^k) - P(x^*) \leq \left(\frac{162\gamma R^2}{\mu k} + \frac{18MR_x}{\mu k^2} + \frac{4}{k^3}\right) \cdot (P(x^0) - P(x^*))$$

If we let

$$\frac{162\gamma R^2}{\mu k} \leq \frac{1}{6}, \quad \frac{18MR_x}{\mu k^2} \leq \frac{1}{6}, \quad \text{and} \quad \frac{4}{k^3} \leq \frac{1}{6},$$

which is equivalent to

$$k \geq O\left(\frac{\gamma R^2}{\mu} + \sqrt{\frac{MR_x}{\mu}} + 1\right),$$

then we can obtain $P(x^k) - P(x^*) \leq \frac{1}{2}(P(x^0) - P(x^*))$. This means that we can reduce the residual $P(x^k) - P(x^*)$ by half after $O\left(\frac{\gamma R^2}{\mu} + \sqrt{\frac{MR_x}{\mu}} + 1\right)$ iterations. Thus, we have $P(x^k) - P(x^*) \leq \epsilon$ as long as

$$k = O\left(\left(\frac{\gamma R^2}{\mu} + \sqrt{\frac{MR_x}{\mu}} + 1\right) \log\left(\frac{P(x^0) - P(x^*)}{\epsilon}\right)\right).$$

C.5 Proof of Theorem 4.5

Since $s^k = \operatorname{argmin}_{s \in \mathbb{R}^d} T(x^k, s)$, from the optimality condition, we have

$$\nabla f(x^k) + \lambda x^k + (\mathbf{H}^k + \lambda \mathbf{I})s^k + \frac{M}{2}\|s^k\|s^k = 0,$$

which indicates that

$$x^{k+1} = x^k - \left(\mathbf{H}^k + \lambda \mathbf{I} + \frac{M}{2}\|x^{k+1} - x^k\| \right)^{-1} (\nabla f(x^k) + \lambda x^k). \quad (48)$$

Then we can obtain

$$\begin{aligned} & \|x^{k+1} - x^*\| \\ &= \left\| x^k - x^* - \left(\mathbf{H}^k + \lambda \mathbf{I} + \frac{M}{2}\|x^{k+1} - x^k\| \right)^{-1} (\nabla f(x^k) + \lambda x^k) \right\| \\ &= \left\| \left(\mathbf{H}^k + \lambda \mathbf{I} + \frac{M}{2}\|x^{k+1} - x^k\| \right)^{-1} \left(\left(\mathbf{H}^k + \lambda \mathbf{I} + \frac{M}{2}\|x^{k+1} - x^k\| \right) (x^k - x^*) - \nabla f(x^k) - \lambda x^k \right) \right\| \\ &\leq \frac{1}{\mu} \left\| \left(\mathbf{H}^k + \lambda \mathbf{I} + \frac{M}{2}\|x^{k+1} - x^k\| \right) (x^k - x^*) - \nabla f(x^k) - \lambda x^k \right\| \\ &\stackrel{(32)}{=} \frac{1}{\mu} \left\| \mathbf{H}^k(x^k - x^*) - (\nabla f(x^k) - \nabla f(x^*)) + \frac{M}{2}\|x^{k+1} - x^k\|(x^k - x^*) \right\| \\ &\leq \frac{1}{\mu} \left\| \mathbf{H}^k(x^k - x^*) - (\nabla f(x^k) - \nabla f(x^*)) \right\| + \frac{M\|x^{k+1} - x^k\|}{2\mu} \|x^k - x^*\|. \end{aligned}$$

Then same as the analysis in the proof of Theorem 3.5, we can obtain

$$\begin{aligned} \|x^{k+1} - x^*\| &\leq \frac{3\gamma\|x^k - x^*\|}{n\mu} \sum_{i=1}^n \sum_{j=1}^m \frac{\|a_{ij}\|^2}{m} \left| \beta^k - \int_0^1 \frac{h_{ij}(x^* + \tau(x^k - x^*)) + 2\gamma}{h_{ij}^k + 2\gamma} d\tau \right| \\ &\quad + \frac{M\|x^{k+1} - x^k\|}{2\mu} \|x^k - x^*\|. \end{aligned}$$

Combing the above inequality and (41) yields

$$\begin{aligned} & \|x^{k+1} - x^*\|^2 \\ &\leq \left(1 + \frac{1}{3} \right) \frac{9\gamma^2\|x^k - x^*\|^2}{n\mu^2} \sum_{i=1}^n \sum_{j=1}^m \frac{\|a_{ij}\|^4}{m} \left| \beta^k - \int_0^1 \frac{h_{ij}(x^* + \tau(x^k - x^*)) + 2\gamma}{h_{ij}^k + 2\gamma} d\tau \right|^2 \\ &\quad + (1 + 3) \cdot \frac{M^2\|x^{k+1} - x^k\|^2}{4\mu^2} \|x^k - x^*\|^2 \\ &\stackrel{(41)}{\leq} \frac{12\gamma^2\|x^k - x^*\|^2}{n\mu^2} \sum_{i=1}^n \sum_{j=1}^m \frac{\|a_{ij}\|^4}{m} \left(\frac{8\nu^2 R^2}{\gamma^2} \|x^k - x^*\|^2 + \frac{8}{\gamma^2} \max_i \{ \|h_i^k - h_i(x^*)\|^2 \} \right) \\ &\quad + \frac{M^2\|x^{k+1} - x^k\|^2}{\mu^2} \|x^k - x^*\|^2. \end{aligned}$$

Then we can continue and get

$$\begin{aligned}
\|x^{k+1} - x^*\|^2 &\leq \frac{96\|x^k - x^*\|^2 R^4}{\mu^2} \sum_{i=1}^n \|h_i^k - h_i(x^*)\|^2 + \frac{96\nu^2 R^6 \|x^k - x^*\|^4}{\mu^2} \\
&\quad + \frac{M^2 \|x^{k+1} - x^k\|^2}{\mu^2} \|x^k - x^*\|^2 \\
&= \frac{96\|x^k - x^*\|^2 R^4}{\mu^2} \mathcal{H}^k + \frac{96\nu^2 R^6 \|x^k - x^*\|^4}{\mu^2} + \frac{M^2 \|x^{k+1} - x^k\|^2}{\mu^2} \|x^k - x^*\|^2.
\end{aligned}$$

Then from $\|x^{k+1} - x^k\|^2 \leq 2\|x^{k+1} - x^*\|^2 + 2\|x^k - x^*\|^2$, we can get

$$\begin{aligned}
\|x^{k+1} - x^*\|^2 &\leq \frac{96\|x^k - x^*\|^2 R^4}{\mu^2} \mathcal{H}^k + \frac{(96\nu^2 R^6 + 2M^2)\|x^k - x^*\|^4}{\mu^2} \\
&\quad + \frac{2M^2 \|x^{k+1} - x^k\|^2}{\mu^2} \|x^k - x^*\|^2.
\end{aligned} \tag{49}$$

Assume $\|x^k - x^*\|^2 \leq \frac{\mu^2}{432mn\nu^2 R^6}$ for all $k \geq 0$. Then from (49) and $M = \nu R^3$, we have

$$\begin{aligned}
\|x^{k+1} - x^*\|^2 &\leq \frac{\mu^2}{432mn\nu^2 R^6} \cdot \frac{96R^4}{\mu^2} \mathcal{H}^k + \frac{\mu^2}{432mn\nu^2 R^6} \cdot \frac{100\nu^2 R^6 \|x^k - x^*\|^2}{\mu^2} \\
&\leq \frac{2}{9mn\nu^2 R^2} \mathcal{H}^k + \frac{1}{4mn} \|x^k - x^*\|^2,
\end{aligned}$$

and by taking expectation, we have

$$\mathbb{E}\|x^{k+1} - x^*\|^2 \leq \frac{2}{9mn\nu^2 R^2} \mathbb{E}[\mathcal{H}^k] + \frac{1}{4mn} \|x^k - x^*\|^2. \tag{50}$$

Let $\Phi_3^k := \|x^k - x^*\|^2 + \frac{4}{9mn\nu^2 R^2} \mathcal{H}^k$. Combing (50) and the evolution of \mathcal{H}^k in Lemma B.1, we arrive at

$$\begin{aligned}
\mathbb{E}[\Phi_3^{k+1}] &= \mathbb{E}\|x^{k+1} - x^*\|^2 + \frac{4}{9mn\nu^2 R^2} \mathbb{E}[\mathcal{H}^{k+1}] \\
&\leq \frac{4}{9mn\nu^2 R^2} \left(1 - \eta + \frac{\eta}{2}\right) \mathbb{E}[\mathcal{H}^k] + \left(\frac{1}{4mn} + \frac{4}{9}\right) \mathbb{E}\|x^k - x^*\|^2 \\
&\leq \left(1 - \min\left\{\frac{\eta}{2}, \frac{1}{4}\right\}\right) \mathbb{E}[\Phi_1^k],
\end{aligned}$$

which implies that $\mathbb{E}[\Phi_3^k] \leq \left(1 - \min\left\{\frac{\eta}{2}, \frac{1}{4}\right\}\right)^k \Phi_3^0$. Then we further have $\mathbb{E}\|x^k - x^*\|^2 \leq \left(1 - \min\left\{\frac{\eta}{2}, \frac{1}{4}\right\}\right)^k \Phi_3^0$ and $\mathbb{E}[\mathcal{H}^k] \leq \left(1 - \min\left\{\frac{\eta}{2}, \frac{1}{4}\right\}\right)^k 3mn\nu^2 R^2 \Phi_3^0$. Assume $x^k \neq x^*$ for all k . Then from (49), we have

$$\frac{\|x^{k+1} - x^*\|^2}{\|x^k - x^*\|^2} \leq \frac{96R^4}{\mu^2} \mathcal{H}^k + \frac{98\nu^2 R^6}{\mu^2} \|x^k - x^*\|^2 + \frac{2\nu^2 R^6}{\mu^2} \|x^{k+1} - x^*\|^2,$$

and by taking expectation, we can get

$$\begin{aligned}
\mathbb{E}\left[\frac{\|x^{k+1} - x^*\|^2}{\|x^k - x^*\|^2}\right] &\leq \frac{96R^4}{\mu^2} \mathbb{E}[\mathcal{H}^k] + \frac{98\nu^2 R^6}{\mu^2} \mathbb{E}\|x^k - x^*\|^2 + \frac{2\nu^2 R^6}{\mu^2} \mathbb{E}\|x^{k+1} - x^*\|^2 \\
&\leq \left(1 - \min\left\{\frac{\eta}{2}, \frac{1}{4}\right\}\right)^k (3mn\eta + 1) \frac{100\nu^2 R^6}{\mu^2} \Phi_3^0.
\end{aligned}$$

D Extra Method: MAX-NEWTON

In this section we propose and analyze one more method, MAX-NEWTON (MN), which should be seen as a variant of NEWTON-STAR (NS). Like NS, MS is not practical and is of theoretical interest only. This was the first method we developed, and all subsequent development that eventually lead to the results in this paper started here. MS differs from NS in how we approximate Hessian of $P(x)$. As in NS, we also assume that we know all $h_{ij}(x^*)$ at the optimum. However, we estimate the Hessian at x^k by $\mathbf{H}^k = \frac{1}{n} \sum_{i=1}^n \mathbf{H}_i^k$, where \mathbf{H}_i^k is defined differently:

$$\mathbf{H}_i^k = \frac{1}{m} \max_{j \in [m]} \left\{ \frac{h_{ij}(x^k)}{h_{ij}(x^*)} \right\} \sum_{j=1}^m h_{ij}(x^*) a_{ij} a_{ij}^\top = \frac{\beta_i^k}{m} \sum_{j=1}^m h_{ij}(x^*) a_{ij} a_{ij}^\top.$$

Above, we define $\beta_i^k := \max_{j \in [m]} \frac{h_{ij}(x^k)}{h_{ij}(x^*)}$. Subsequently, we perform a Newton-like step:

$$x^{k+1} = x^k - (\mathbf{H}^k + \lambda \mathbf{I})^{-1} \left(\frac{1}{n} \sum_{i=1}^n \nabla f_i(x^k) + \lambda x^k \right).$$

The method is summarized as Algorithm 4.

Algorithm 4 MN: MAX-NEWTON

Initialization: $x^0 \in \mathbb{R}^d$
for $k = 0, 1, 2, \dots$ **do**
 Broadcast x^k to all workers
 for $i = 1, \dots, n$ **do**
 Compute $\nabla f_i(x^k)$
 $\beta_i^k = \max_{j \in [m]} \frac{h_{ij}(x^k)}{h_{ij}(x^*)}$
 Send $\nabla f_i(x^k)$ and β_i^k to the server
 end for
 $\mathbf{H}_i^k = \frac{\beta_i^k}{m} \sum_{j=1}^m h_{ij}(x^*) a_{ij} a_{ij}^\top$
 $\mathbf{H}^k = \frac{1}{n} \sum_{i=1}^n \mathbf{H}_i^k$
 $x^{k+1} = x^k - (\mathbf{H}^k + \lambda \mathbf{I})^{-1} \left(\frac{1}{n} \sum_{i=1}^n \nabla f_i(x^k) + \lambda x^k \right)$
end for

We now show that like original Newton's method and NEWTON-STAR, MAX-NEWTON also converges locally quadratically.

Theorem D.1 (Local quadratic convergence). *Assume that f_{ij} is convex for all i, j and that we know $h_{ij}(x^*)$ for all i, j . Furthermore, assume that $\mathbf{H}_{ij}(x^*) \succeq \mu^* \mathbf{I}$ for some $\mu^* > 0$ for all i, j (for instance, this holds if f_{ij} is μ^* -strongly convex). Then for any starting point $x^0 \in \mathbb{R}^d$, the iterates of MAX-NEWTON for solving problem (1) satisfy the following inequality*

$$\|x^{k+1} - x^*\| \leq \left(\frac{\nu}{2\lambda} \frac{1}{nm} \sum_{i=1}^n \sum_{j=1}^m \|a_{ij}\|^3 \left[\frac{h_{ij}(x^*) R}{\mu^* \|a_{ij}\|} + 1 \right] \right) \|x^k - x^*\|^2,$$

where $R := \max_{ij} \|a_{ij}\|$.

Proof. Since $\mathbf{H}_{ij}(x^*) \succeq \mu^* \mathbf{I}$, then $h_{ij}(x^*)$ is positive for all i, j . We estimate Hessian at x^k by $\mathbf{H}^k = \frac{1}{n} \sum_{i=1}^n \mathbf{H}_i^k$, where

$$\mathbf{H}_i^k = \frac{1}{m} \max_{j \in [m]} \left\{ \frac{h_{ij}(x^k)}{h_{ij}(x^*)} \right\} \sum_{j=1}^m h_{ij}(x^*) a_{ij} a_{ij}^\top = \frac{\beta_i^k}{m} \sum_{j=1}^m h_{ij}(x^*) a_{ij} a_{ij}^\top.$$

Notice that $h_{ij}(x) \geq 0$, that's why $\beta_i^k \geq 0$ too. In consequence, $\mathbf{H}_i^k \succeq \mathbf{0}$, and $\mathbf{H}^k + \lambda \mathbf{I} \succeq \lambda \mathbf{I}$. Then we have

$$\begin{aligned} \|x^{k+1} - x^*\| &= \left\| x^k - x^* (\mathbf{H}^k + \lambda \mathbf{I})^{-1} (\nabla f(x^k) + \lambda x^k) \right\| \\ &= \left\| (\mathbf{H}^k + \lambda \mathbf{I})^{-1} ((\mathbf{B}^k + \lambda \mathbf{I})(x^k - x^*) - \nabla f(x^k) - \lambda x^k) \right\| \\ &\leq \left\| (\mathbf{H}^k + \lambda \mathbf{I})^{-1} \right\| \left\| (\mathbf{H}^k + \lambda \mathbf{I})(x^k - x^*) - \nabla f(x^k) - \lambda x^k \right\| \\ &\leq \frac{1}{\lambda} \left\| (\mathbf{H}^k + \lambda \mathbf{I})(x^k - x^*) - \nabla f(x^k) - \lambda x^k \right\| \\ &\stackrel{(32)}{=} \frac{1}{\lambda} \left\| \mathbf{H}^k(x^k - x^*) - (\nabla f(x^k) - \nabla f(x^*)) \right\| \\ &\leq \frac{1}{n\lambda} \sum_{i=1}^n \left\| \mathbf{H}_i^k(x^k - x^*) - (\nabla f_i(x^k) - \nabla f_i(x^*)) \right\| \\ &\leq \frac{1}{nm\lambda} \sum_{i=1}^n \sum_{j=1}^m \left\| \beta_i^k h_{ij}(x^*) a_{ij} a_{ij}^\top (x^k - x^*) \right. \\ &\quad \left. - \int_0^1 h_{ij}(x^* + \tau(x^k - x^*)) a_{ij} a_{ij}^\top (x^k - x^*) d\tau \right\|. \end{aligned}$$

Finally we obtain

$$\|x^{k+1} - x^*\| \leq \frac{\|x^k - x^*\|}{nm\lambda} \sum_{i=1}^n \sum_{j=1}^m \|a_{ij}\|^2 h_{ij}(x^*) \left| \beta_i^k - \int_0^1 \frac{h_{ij}(x^* + \tau(x^k - x^*))}{h_{ij}(x^*)} d\tau \right|. \quad (51)$$

Now we want to upper bound the last term. We know that $h_{ij}(x)$ is $\nu \|a_{ij}\|$ -Lipschitz function, then

$$\begin{aligned} h_{ij}(x^*) - \nu \|a_{ij}\| \tau \|x^k - x^*\| &\leq h_{ij}(x^* + \tau(x^k - x^*)) \leq h_{ij}(x^*) + \nu \|a_{ij}\| \tau \|x^k - x^*\| \\ 1 - \frac{\nu \|a_{ij}\|}{h_{ij}(x^*)} \tau \|x^k - x^*\| &\leq \frac{h_{ij}(x^* + \tau(x^k - x^*))}{h_{ij}(x^*)} \leq 1 + \frac{\nu \|a_{ij}\|}{h_{ij}(x^*)} \tau \|x^k - x^*\|. \end{aligned}$$

We integrate these two inequalities and obtain

$$1 - \frac{\nu \|a_{ij}\|}{2h_{ij}(x^*)} \|x^k - x^*\| \leq \int_0^1 \frac{h_{ij}(x^* + \tau(x^k - x^*))}{h_{ij}(x^*)} d\tau \leq 1 + \frac{\nu \|a_{ij}\|}{2h_{ij}(x^*)} \|x^k - x^*\|.$$

Let's denote $j_i^k = \arg \max_{j \in [m]} \left\{ \frac{h_{ij}(x^k)}{h_{ij}(x^*)} \right\}$, then $\beta_i^k = \frac{h_{ij_i^k}(x^k)}{h_{ij_i^k}(x^*)}$. It means that for β_i^k we can write the same two inequalities as above

$$1 - \frac{\nu \|a_{ij_i^k}\|}{2h_{ij_i^k}(x^*)} \|x^k - x^*\| \leq \beta_i^k \leq 1 + \frac{\nu \|a_{ij_i^k}\|}{2h_{ij_i^k}(x^*)} \|x^k - x^*\|.$$

Taking minimum and maximum in two parts we get

$$1 - \frac{\nu \max_{ij} \|a_{ij}\|}{2 \min_{ij} h_{ij}(x^*)} \|x^k - x^*\| \leq \beta_i^k \leq 1 + \frac{\nu \max_{ij} \|a_{ij}\|}{2 \min_{ij} h_{ij}(x^*)} \|x^k - x^*\|.$$

Using inequalities for β_i^k and for the integral we obtain

$$|\beta_i^k - h_{ij}(x^k)| \leq \frac{\nu \|x^k - x^*\|}{2} \left[\frac{\max_{ij} \|a_{ij}\|}{\min_{ij} h_{ij}(x^*)} + \frac{\|a_{ij}\|}{h_{ij}(x^*)} \right].$$

Combining above and (51) we get

$$\begin{aligned} \|x^{k+1} - x^*\| &\leq \frac{\nu \|x^k - x^*\|^2}{2nm\lambda} \sum_{i=1}^n \sum_{j=1}^m \|a_{ij}\|^2 h_{ij}(x^*) \left[\frac{\max_{ij} \|a_{ij}\|}{\min_{ij} h_{ij}(x^*)} + \frac{\|a_{ij}\|}{h_{ij}(x^*)} \right] \\ &\leq \frac{\nu \|x^k - x^*\|^2}{2nm\lambda} \sum_{i=1}^n \sum_{j=1}^m \|a_{ij}\|^3 \left[\frac{h_{ij}(x^*) \max_{ij} \|a_{ij}\|}{\mu^* \|a_{ij}\|} + 1 \right], \end{aligned}$$

where we use inequality $h_{ij}(x^*) \geq \mu^*$. □

As we can see, the size of convergence area for **NS** is larger than for **MN**, because **NS** depends on $\|a_{ij}\|^3$, while **MN** depends on $\|a_{ij}\|^2 \max_{i,j} \|a_{ij}\|$.

References

- [1] D. Alistarh, D. Grubic, J. Li, R. Tomioka, and M. Vojnovic. QSGD: Communication-efficient SGD via gradient quantization and encoding. *Advances in Neural Information Processing Systems*, pages 1709–1720, 2017.
- [2] Amir Beck. *Introduction to Nonlinear Optimization: Theory, Algorithms, and Applications with MATLAB*. Society for Industrial and Applied Mathematics, USA, 2014. ISBN 1611973643.
- [3] Ron Bekkerman, Mikhail Bilenko, and John Langford. *Scaling up machine learning: Parallel and distributed approaches*. Cambridge University Press, 2011.
- [4] J. Bernstein, Y. X. Wang, K. Azizzadenesheli, and A. Anandkumar. SignSGD: Compressed optimisation for non-convex problems. *The 35th International Conference on Machine Learning*, pages 560–569, 2018.
- [5] Aleksandr Beznosikov, Samuel Horváth, Peter Richtárik, and Mher Safaryan. On biased compression for distributed learning. *arXiv:2002.12410*, 2020.
- [6] Charles G Broyden. Quasi-newton methods and their application to function minimisation. *Mathematics of Computation*, 21(99):368–381, 1967.
- [7] Rixon Crane and Fred Roosta. DINGO: Distributed Newton-type method for gradient-norm optimization. In H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché-Buc, E. Fox, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 32, pages 9498–9508. Curran Associates, Inc., 2019. URL <https://proceedings.neurips.cc/paper/2019/file/9718db12cae6be37f7349779007ee589-Paper.pdf>.
- [8] Aaron Defazio, Francis Bach, and Simon Lacoste-Julien. SAGA: A fast incremental gradient method with support for non-strongly convex composite objectives. In Z. Ghahramani, M. Welling, C. Cortes, N. D. Lawrence, and K. Q. Weinberger, editors, *Advances in Neural Information Processing Systems 27*, pages 1646–1654. Curran Associates, Inc., 2014.
- [9] Rodger Fletcher. A new approach to variable metric algorithms. *The Computer Journal*, 13(3):317–323, 1970.
- [10] Avishek Ghosh, Raj Kumar Maity, Arya Mazumdar, and Kannan Ramchandran. Communication efficient distributed approximate Newton method. In *IEEE International Symposium on Information Theory (ISIT)*, 2020. doi: 10.1109/ISIT44484.2020.9174216.
- [11] Donald Goldfarb. A family of variable-metric methods derived by variational means. *Mathematics of Computation*, 24(109):23–26, 1970.
- [12] Nicholas IM Gould, Daniel P Robinson, and H Sue Thorne. On solving trust-region and other regularised subproblems in optimization. *Mathematical Programming Computation*, 2(1):21–57, 2010.

- [13] Andreas Griewank. The modification of Newton’s method for unconstrained optimization by bounding cubic terms. Technical report, Department of Applied Mathematics and Theoretical Physics, University of Cambridge, 1981. Technical Report NA/12.
- [14] Filip Hanzely, Nikita Doikov, Yurii Nesterov, and Peter Richtárik. Stochastic subspace cubic Newton method. In *International Conference on Machine Learning*, pages 4027–4038. PMLR, 2020.
- [15] Samuel Horváth, Chen-Yu Ho, Ľudovít Horváth, Atal Narayan Sahu, Marco Canini, and Peter Richtárik. Natural compression for distributed deep learning. *arXiv preprint arXiv:1905.10988*, 2019.
- [16] Samuel Horváth, Dmitry Kovalev, Konstantin Mishchenko, Sebastian Stich, and Peter Richtárik. Stochastic distributed learning with gradient quantization and variance reduction. *arXiv preprint arXiv:1904.05115*, 2019.
- [17] Rustem Islamov, Xun Qian, and Peter Richtárik. Distributed second order methods with fast rates and compressed communication. *arXiv preprint arXiv:2102.07158*, 2021.
- [18] Rie Johnson and Tong Zhang. Accelerating stochastic gradient descent using predictive variance reduction. In *NIPS*, pages 315–323, 2013.
- [19] Ahmed Khaled, Konstantin Mishchenko, and Peter Richtárik. Tighter theory for local SGD on identical and heterogeneous data. In *The 23rd International Conference on Artificial Intelligence and Statistics (AISTATS 2020)*, 2020.
- [20] Sarit Khirirat, Hamid Reza Feyzmahdavian, and Mikael Johansson. Distributed learning with compressed gradients. In *arXiv preprint arXiv:1806.06573*, 2018.
- [21] Dmitry Kovalev, Konstantin Mishchenko, and Peter Richtárik. Stochastic Newton and cubic Newton methods with simple local linear-quadratic rates. In *NeurIPS Beyond First Order Methods Workshop*, 2019.
- [22] Dmitry Kovalev, Robert M. Gower, Peter Richtárik, and Alexander Rogozin. Fast linear convergence of randomized BFGS. *arXiv:2002.11337*, 2020.
- [23] Zhize Li, Dmitry Kovalev, Xun Qian, and Peter Richtárik. Acceleration for compressed gradient descent in distributed and federated optimization. In *International Conference on Machine Learning*, 2020.
- [24] Hongzhou Lin, Julien Mairal, and Zaid Harchaoui. A universal catalyst for first-order optimization. *arXiv preprint arXiv:1506.02186*, 2015.
- [25] Nicolas Loizou and Peter Richtárik. Linearly convergent stochastic heavy ball method for minimizing generalization error. In *NIPS Workshop on Optimization for Machine Learning*, 2017.
- [26] Nicolas Loizou and Peter Richtárik. Momentum and stochastic momentum for stochastic gradient, Newton, proximal point and subspace descent methods. *arXiv:1712.09677*, 2017.

- [27] Chenxin Ma, Jakub Konečný, Martin Jaggi, Virginia Smith, Michael I. Jordan, Peter Richtárik, and Martin Takáč. Distributed optimization with arbitrary local solvers. *Optimization Methods and Software*, 32(4):813–848, 2017.
- [28] Yura Malitsky and Konstantin Mishchenko. Adaptive gradient descent without descent. In Hal Daumé III and Aarti Singh, editors, *Proceedings of the 37th International Conference on Machine Learning*, volume 119 of *Proceedings of Machine Learning Research*, pages 6702–6712. PMLR, 13–18 Jul 2019.
- [29] Konstantin Mishchenko, Eduard Gorbunov, Martin Takáč, and Peter Richtárik. Distributed learning with compressed gradient differences. *arXiv preprint arXiv:1901.09269*, 2019.
- [30] Konstantin Mishchenko, Ahmed Khaled, and Peter Richtárik. Random reshuffling: Simple analysis with vast improvements. In *34th Conference on Neural Information Processing Systems (NeurIPS 2020)*, 2020.
- [31] Deanna Needell, Nathan Srebro, and Rachel Ward. Stochastic gradient descent, weighted sampling, and the randomized Kaczmarz algorithm. *Mathematical Programming*, 155(1–2):549–573, 2015.
- [32] Arkadi Nemirovski, Anatoli Juditsky, Guanghui Lan, and Alexander Shapiro. Robust stochastic approximation approach to stochastic programming. *SIAM Journal on Optimization*, 19(4):1574–1609, 2009.
- [33] Yurii Nesterov and Boris T. Polyak. Cubic regularization of Newton method and its global performance. *Mathematical Programming*, 108(1):177–205, 2006.
- [34] Boris Polyak and Andrey Tremba. New versions of Newton method: step-size choice, convergence domain and under-determined equations. *arXiv preprint arXiv:1703.07810*, 2019.
- [35] Boris Polyak and Andrey Tremba. New versions of newton method: step-size choice, convergence domain and under-determined equations. *Optimization Methods and Software*, 35(6):1272–1303, 2020.
- [36] Joseph Raphson. Analysis aequationum universalis seu ad aequationes algebraicas resolvendas methodus generalis, & expedita, ex nova infinitarum serierum methodo, deducta ac demonstrata. *Oxford: Richard Davis*, 1697.
- [37] Sashank J. Reddi, Jakub Konečný, Peter Richtárik, Barnabás Póczos, and Alex Smola. AIDE: fast and communication efficient distributed optimization. *arXiv:1608.06879*, 2016.
- [38] H. Robbins and S. Monro. A stochastic approximation method. *Annals of Mathematical Statistics*, 22:400–407, 1951.
- [39] M. Schmidt, N. Le Roux, and F. Bach. Minimizing finite sums with the stochastic average gradient. *Math. Program.*, 162(1-2):83–112, 2017.

- [40] F. Seide, H. Fu, J. Droppo, G. Li, and D. Yu. 1-bit stochastic gradient descent and its application to data- parallel distributed training of speech DNNs. *Fifteenth Annual Conference of the International Speech Communication Association*, 2014.
- [41] Shai Shalev-Shwartz and Shai Ben-David. *Understanding machine learning: from theory to algorithms*. Cambridge University Press, 2014.
- [42] Ohad Shamir, Nati Srebro, and Tong Zhang. Communication-efficient distributed optimization using an approximate Newton-type method. In *Proceedings of the 31st International Conference on Machine Learning, PMLR*, volume 32, pages 1000–1008, 2014.
- [43] David F Shanno. Conditioning of quasi-Newton methods for function minimization. *Mathematics of computation*, 24(111):647–656, 1970.
- [44] S. U. Stich and S. P. Karimireddy. The error-feedback framework: Better rates for SGD with delayed gradients and compressed communication. *arXiv: 1909.05350*, 2019.
- [45] Sebastian U. Stich. Local SGD converges fast and communicates little. In *International Conference on Learning Representations*, 2020.
- [46] Sebastian U. Stich, Jean-Baptiste Cordonnier, and Martin Jaggi. Sparsified SGD with memory. In *Advances in Neural Information Processing Systems*, pages 4447–4458, 2018.
- [47] Martin Takáč, Avleen Bijral, Peter Richtárik, and Nathan Srebro. Mini-batch primal and dual methods for SVMs. In *30th International Conference on Machine Learning*, pages 537–552, 2013.
- [48] H. Tang, X. Lian, T. Zhang, and J. Liu. DoubleSqueeze: Parallel stochastic gradient descent with double-pass error-compensated compression. In *Proceedings of the 36th International Conference on Machine Learning*, pages 6155–6165, 2019.
- [49] John Wallis. A treatise of algebra, both historical and practical. *Philosophical Transactions of the Royal Society of London*, 15(173):1095–1106, 1685. doi: 10.1098/rstl.1685.0053. URL <https://royalsocietypublishing.org/doi/abs/10.1098/rstl.1685.0053>.
- [50] Shusen Wang, Fred Roosta, Peng Xu, and Michael W Mahoney. Giant: Globally improved approximate newton method for distributed optimization. In S. Bengio, H. Wallach, H. Larochelle, K. Grauman, N. Cesa-Bianchi, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 31, pages 2332–2342. Curran Associates, Inc., 2018. URL <https://proceedings.neurips.cc/paper/2018/file/dabd8d2ce74e782c65a973ef76fd540b-Paper.pdf>.
- [51] W. Wen, C. Xu, F. Yan, C. Wu, Y. Wang, and H. Li. Terngrad: Ternary gradients to reduce communication in distributed deep learning. *Advances in Neural Information Processing Systems*, pages 1509–1519, 2017.

- [52] Lin Xiao and Tong Zhang. A proximal stochastic gradient method with progressive variance reduction. *SIAM Journal on Optimization*, 24(4):2057–2075, 2014.
- [53] Jiaqi Zhang, Keyou You, and Tamer Başar. Distributed adaptive Newton methods with globally superlinear convergence. *arXiv preprint arXiv:2002.07378*, 2020.
- [54] Yuchen Zhang and Lin Xiao. DiSCO: Distributed optimization for self-concordant empirical loss. In *Proceedings of the 32nd International Conference on Machine Learning*, PMLR, volume 37, pages 362–370, 2015.
- [55] Peilin Zhao and Tong Zhang. Stochastic optimization with importance sampling. *The 32nd International Conference on Machine Learning*, 37:1–9, 2015.